

Heroes, Villains, and Victims: Character Narratives in the WPS agenda of the UNSC

Raluca Rilla
Universität Potsdam

Hannah Steinbach
Universität Potsdam

Imge Yüzüncüoğlu
Universität Potsdam

{ralucarilla, hannahsteinbach0312, i.yuezuencue}@gmail.com

Abstract

This paper investigates the narrative construction of character roles (heroes, victims, and villains) in open debates on the United Nations Security Council’s Women, Peace, and Security (WPS) agenda. Drawing from 2,566 speeches delivered between 2000 and 2019, we examine how (gendered) entities are framed within diplomatic discourse. Through a combination of topic modeling, clustering, and supervised learning, we uncover patterns in how women are simultaneously portrayed as empowered agents and vulnerable subjects, while perpetrators are often left unnamed. We manually annotate 54 speeches with role labels and evaluate the effectiveness of multiple computational models, including a fine-tuned RoBERTa classifier and two chat-optimized Large Language Models (DeepSeek-R1 and Llama3.3 70B), in identifying these roles. Our findings reveal substantial variation in model performance, with RoBERTa outperforming LLMs in overall F1 scores, particularly for less frequent and overlapping roles. All code and annotated datasets are publicly available to encourage further research on narrative framing and role attribution in this domain.

1 Introduction

According to [Gehring and Grigoletto \(2023\)](#), "narratives are a crucial group-based mechanism that influences human decision-making" (p. 1). The strategic use of different framing techniques in discourse is particularly salient in the context of open, multilateral debates that represent opportunities for states to construct legitimacy and to influence international norms. In framing different entities as heroes, villains, or victims, authors and speakers subtly shape perceptions of responsibility and agency.

One domain where these narrative strategies play a crucial role is the Women, Peace, and Security (WPS) agenda of the UN Security Council. Ini-

tiated by Resolution 1325 (2000) and adopted on 31 October 2000, the agenda seeks to reshape the role of women in peace and security efforts, making them central figures in discussions on conflict resolution and post-conflict reconstruction. This resolution

"reaffirms the important role of women in the conflict prevention and resolution, peace negotiations, peace-building, peacekeeping, humanitarian response and post-conflict reconstruction" ([UNDP, 2023](#), p. 5).

Subsequent resolutions have expanded and refined these objectives.

The WPS agenda is structured around four pillars: prevention, participation, protection, relief, and recovery. Its overarching goals can be summarized as follows: preventing violence against women and girls in conflict and post-conflict situations, implementing measures to enhance gender equality, increasing women’s participation in all United Nations peace and security efforts, and integrating a gender perspective ([UNDP, 2023](#)).

Given these objectives alone, one can already anticipate that women are positioned as central figures in these debates, yet they are also framed in specific roles. On the one hand, women are framed as victims, suffering in armed conflict and requiring protection. On the other hand, they are also portrayed as heroes, as their (potential) participation in UN peace and security efforts is framed as valuable and necessary. However, the very call to increase their involvement implies that they are still largely excluded from these processes, reinforcing their victimhood. This dual framing creates a complex picture: women are victims because they are denied access to decision-making, yet they are also heroes because their engagement is seen as beneficial. These contrasting framings are reflected not

only in lexical choices (e. g., "suffering" vs. "participation") but also in sentence structure and agency. When depicted as victims, women are passive subjects—violence is inflicted upon them, and they lack control over their circumstances. In contrast, when framed as heroes, their agency is acknowledged, but they are still positioned as dependent on external support. The UN itself is framed as a hero, both for working to increase women’s participation and for its efforts to protect them from violence in conflict zones. Since the objectives are a key component of the agenda, we can expect the sessions themselves to be held in a similar narrative.

This study aims to analyze how heroes, victims, and villains are depicted in open debates on the WPS agenda and to evaluate the ability of large language models (LLMs) to accurately identify these roles¹. The paper is structured as follows: First, we present results from an exploratory meta-analysis of gender representation in the Security Council. Next, we use clusterings and BERTopic (Grootendorst, 2022) to examine the topics discussed in WPS debates and how gendered entities are integrated into these topics. We then refine this analysis using supervised learning, fine-tuning a RoBERTa model to classify the roles of heroes, victims and villains automatically. Finally, we compare two chat-optimized LLMs, Llama and DeepSeek, through prompting.

Contributions of this work are therefore:

1. A meta-analysis providing insights into gender representation in the UN Security Council.
2. An analysis of the most frequent topics in the WPS agenda using clustering, alongside a focused comparison of how gendered entities ("women" and "men") are discussed.
3. A corpus annotated with hero, victim, and villain roles consisting of 54 speeches, as well as annotation guidelines.
4. Results from multiple computational approaches for automatically identifying these roles, including clustering, fine-tuned RoBERTa classification, and prompting with Llama and DeepSeek.

All datasets and used codes are public available on [GitHub](#).

¹The specific definitions used to characterize these roles are detailed in the Annotation Guidelines in Section [Appendix C](#)

2 Related Work

According to Gibbings (2011), "operating at the UN is akin to acquiring a second language" (p. 533). Gibbings, who worked for WILPF-UN and the NGO Working Group on Women, Peace and Security, returned to the UN headquarters for her Master’s thesis to conduct four months of anthropological fieldwork, focusing on UN Security Resolution 1325. This quote refers to the cultural norms at the UN that speakers must adhere to, one of which is the requirement to frame issues in a positive light: "UN speech styles encourage positive visions and utopian dreams" (Gibbings, 2011, p. 534). Furthermore, as discussed in the introduction, she observes that Resolution 1325 is surrounded by discourses that emphasize the knowledge produced by women and position them as peacemakers. Thus, framings that do not align with this portrayal are discouraged.

She illustrates this with the example of two Iraqi women, Amal Al-Khedairy and Nermin Al-Mufti, who were invited to an informal meeting attended by gender officers from UN agencies, NGOs, and several representatives from Member States in October 2003. They were not allowed to formally speak in an open debate of the Security Council, since this was considered too controversial at that time. During this meeting, these women condemned the USA and UK’s invasion as imperialist and criticized the UN for its lack of support, leaving many participants of the meeting "disappointed and embarrassed" (Gibbings, 2011, p. 534). The women were subsequently labeled as "angry". This incident illustrates not only the existence of a specific etiquette and framing within the UN, especially in the context of the WPS agenda and women, but also that women themselves are "expected to embody and reproduce this image as upholders of peace in their communities" (Gibbings, 2011, p. 534).

The UN’s tendency to depict women primarily as victims or 'agents of change' has been widely criticized. Başer (2024), for instance, warns that these framings only reinforce gender norms by failing to acknowledge women also as perpetrators of violence themselves and by excluding men and LGBTQ+ individuals from victim portrayal. Thus,

"the WPS agenda’s reinforcement of peaceful or victim women stereotypes may inadvertently reinforce the notion that men are inherently violent or not

worthy of protection." (Başer, 2024, p.15)

This view is supported by Carpenter (2005), who finds that the UN's civilian protection discourse is deeply gendered, frequently invoking the phrase "women and children" to refer to civilians while excluding men. This framing emphasizes women's and children's perceived vulnerability, also serving to attract donor aid and international cooperation (Carpenter, 2005).

Nevertheless, the discourse is constantly evolving. de la Rosa and Lázaro (2019), for example, analyze the portrayal of women in the eight WPS resolutions adopted between Resolution 1325 in 2000 and Resolution 2242 in 2015, focusing on metaphor usage. They find that the framing of women has shifted from victimhood to heroism:

"the narrative in resolutions changes from a perspective where women need to be protected from sexual violence to another in which they need to be empowered and involved in the fight against it" (de la Rosa and Lázaro, 2019, p.382).

This suggests that, at least within the resolutions, the representation of women in the WPS agenda has evolved.

However, since open debates allow various Member States to express their views—unlike resolutions, which are negotiated and largely agreed upon—the focus here is on how different entities are portrayed in these debates. Although their research does not explicitly address victim, villain, and hero framings, True and Wiener (2019) provide valuable insights into these characterizations. Analyzing the 15th-anniversary Security Council open debate on WPS in 2015, they highlight the agenda's strong focus on sexual violence and echo the concerns of WPS agenda founders that

"contemporary social validation of that agenda based on the sexual violence perpetrated against women and girls runs the risk of reinforcing women's and girls' status as victims in respect of peace and security" (True and Wiener, 2019, p. 565).

Based on these findings and previous research, we hypothesize that women and girls are predominantly portrayed as victims in open debates.

Regarding heroism, True and Wiener (2019) identify a group of Member States—mainly from the Global North—that strongly support the WPS agenda and often emphasize their own achievements, thus presenting themselves as heroes. A similar pattern emerges in another group of states affected by conflict. While these states endorse the WPS agenda, they are constrained in their ability to implement it independently and require external support. As a result, they position themselves both as heroes (for their efforts) and as victims (for the challenges they face). True and Wiener (2019) also mention two other groups: states that seek to limit the scope of the agenda, citing national sovereignty (e. g., Russia, China, and Egypt), and those that support the agenda without contesting its expansion but still insist on their autonomy. It remains to be seen whether these latter two groups also emphasize their achievements in a heroic light as much as the first two.

Finally, True and Wiener (2019) note that the portrayal of perpetrators can be particularly contentious. While states are largely united in condemning violence by non-state actors such as ISIS or Boko Haram, directly calling out fellow Member States proves more challenging:

"It is easier to object to breaches of the norm by non-state actors not present in the open debate, however, than to name and shame breaches made by states present" (True and Wiener, 2019, p. 563).

This review has explored how different entities are framed within the WPS discourse, with a particular emphasis on the central figures—women. While initial portrayals primarily focused on women's victimhood, analyses suggest a shift toward framing them as 'agents of peace' with valuable knowledge and a strong commitment to peace, thus positioning them as heroes. However, these representations remain controversial, as they obscure women's roles as perpetrators of violence and often downplay the victimhood of men, sometimes for strategic reasons such as securing funding.

True and Wiener (2019)'s analysis of an open debate has provided initial insights into these character framings, revealing the victimization of women (and, in some cases, states), self-portrayals of heroism by Member States, and the targeting of non-state actors as perpetrators, while avoiding direct criticism of present states. Given the UN's

tendency to maintain a positive tone, we expect an emphasis on achievements and successes—thus, on heroes. Our next step is to assess whether these insights hold across a broader set of open debates and to explore the potential of machine learning approaches to process large volumes of text, enabling more efficient analysis and comparison. The following section will discuss existing approaches for automatic character framing analysis.

NLP Approaches

The identification of roles such as heroes, victims and villains is not a recent development. In the field of Natural Language Processing (NLP), the identification of such roles is of considerable importance, particularly for the detection of framing strategies. As noted by [Gomez-Zara et al. \(2018a\)](#), framing is a widely used tool in the media to shape audience perception by implicitly assigning roles to entities. To counter this influence, [Gomez-Zara et al. \(2018a\)](#) developed a browser extension that identifies and highlights how entities are framed (as heroes, victims or villains) while reading news articles online.

The system works in several steps. First, named entity recognition (NER) is used to identify key figures in the text. These entities are then ranked in order of importance based on factors such as frequency of mention. Next, the system analyzes the role framing of each entity by examining co-occurring terms and assessing their similarity to generic definitions of heroes, victims and villains. Sentiment scores are then calculated for the associated terms: positive or neutral terms are associated with heroes, while negative or neutral terms are associated with victims or villains. Finally, the predicted character roles are assigned to the main entities and displayed to the user within the browser interface.

The implementation has not undergone a formal evaluation yet; however, it is acknowledged that the browser extension may struggle with complex sentence structures, which could affect its accuracy ([Gomez-Zara et al., 2018a](#)).

[Stammbach et al. \(2022\)](#) argue that large pre-trained language models such as GPT-3 are more effective at labeling character roles than previously used NLP methods. Their approach frames the task as a reading comprehension or closed-ended question-answer task, where GPT-3 is given both the relevant question (e. g., "Who is the

hero/villain?") and the context.

To evaluate their method, [Stammbach et al. \(2022\)](#) conduct multiple experiments across different text domains, rather than limiting their analysis to a single genre. They apply their approach to newspaper articles about fracking, plot summaries of Disney movies, and a corpus of US State of the Union addresses, the latter being an unlabelled dataset. Their results show that GPT-3 achieves higher accuracy than [Gomez-Zara et al. \(2018a\)](#) in the first experiment. The model performs particularly well on Disney film plots, effectively identifying character roles. In addition, their third experiment, which uses a social science-oriented evaluation, further supports the effectiveness of their approach. These results show that LLMs not only outperform previous methods in role labelling, but are also easier to implement.

While media, such as news reports, often use stereotypes to quickly associate entities with specific roles, [Sap et al. \(2017\)](#) focus on investigating how these stereotypes are formed in the first place. To this end, they extend the connotation frame lexicon developed by [Rashkin et al. \(2015\)](#) by introducing a verb list specifically designed to capture 'power' and 'agency'. Verbs play a crucial role in conveying who acts actively or passively, who has dominance, and how authority is distributed within a narrative.

In their framework, "power" reflects the degree of authority expressed by a verb, while "agency" measures the extent to which a character is "is powerful, decisive, and capable of pushing forward their own storyline" ([Sap et al., 2017](#), p.4).

To empirically test their framework, [Sap et al. \(2017\)](#) conduct a quantitative study analyzing gender bias in modern films. Using their verb lists, they compare the portrayal of male and female characters, assessing differences in how agency and power are attributed to each gender. Their approach provides a methodological framework for analyzing narrative dynamics in different contexts, making it a valuable tool for studying power structures in media representation.

A key aspect of developing approaches for automatic role detection is not only the selection of the appropriate method, but also the definition of the dataset and its annotation process. [Frermann et al. \(2023\)](#) criticize existing datasets and frameworks for assuming that entities can only be assigned a single role, rather than recognizing that entities can be associated with multiple roles simultaneously.

To address this limitation, they introduce an annotation process in which annotators answer a set of predefined questions to determine an entity’s role. Specifically, annotators assess whether an article contains an entity that

"could *alleviate the problem* (Hero); was *responsible for the problem* (Villain); or was *negatively affected by the problem* (Victim)" (Frermann et al., 2023, p. 4).

If an annotator answers with "yes" to any of these questions, they are then asked to identify the most appropriate entity within the text.

This question-based annotation framework allows entities to be assigned multiple overlapping roles, facilitating a multi-label classification approach. By integrating this method, Frermann et al. (2023) contributes to a fine-grained understanding of narrative roles in discourse, overcoming the limitations of single-label classification.

As we have seen that LLMs appear to outperform previously used NLP methods, this project uses three LLMs—RoBERTa, DeepSeek, and Llama3—to investigate and compare how characters and their associated labels can be identified in the context of UN Security Council debates on WPS. However, it is important to note that this is only one aspect of our broader analysis. The primary goal of our research is to explore how entities within these debates are characterized according to labels such as ‘hero’, ‘victim’, and ‘villain’. The use of LLMs is specifically tested to assess whether they can be employed for future analyses of larger datasets, enabling more meaningful and accurate statements on how entities are portrayed across different time periods or specific countries. Through this approach, we aim to contribute to the automation of discourse analysis, allowing for more efficient and scalable studies of international relations and gendered narratives.

3 Dataset and Task

3.1 Data

Using Schoenfeld et al. (2019)’s dataset of all UNSC debates between 1997 and 2019 (inclusive), we manually extracted the sessions addressing the WPS agenda². The original dataset includes 82165 speeches across 5748 debates. Of

these 5748, 68 sessions (4613 speeches, including those of the President and the Secretary-General) were centered on the topic of women, peace, and security. We identified and removed the duplicate of a debate, which appeared under two distinct file beginnings: UNSC_2010_SPV.6393 and UNSC_2010_SPV.6396Resumption1. Additionally, the dataset (as well as the official Security Council source we used for the identification of the WPS debates) erroneously included the proceedings of S/PV.6180 from 2009 under the meeting of the year 2001. Upon further investigation, we concluded that there was no meeting of the UNSC that discussed items on the WPS agenda in 2001.

3.2 Data Cleaning

As an initial preprocessing step, we excluded all presidential speeches, as they mainly consist of formalities, such as introducing speakers. After removing those, 2566 speeches remained for our dataset. From the remaining speeches, we extracted information on the year, session, speaker’s country or organization, speaker’s gender (when indicated by titles such as Mr. or Ms.), and the language of the speech, using regular expressions. These metadata were then used in subsequent experiments.

3.3 Manual Annotation of Dataset

In trying to establish our annotation guidelines, we first randomly extracted four speeches from the dataset to individually annotate heroes, victims, and villains using generic definitions. In discussing our reasoning for certain role assignment choices, we began elaborating a set of rules and definitions specific to our corpus and goal of understanding the portrayal of gender in these speeches. Notably, as none of us annotators are native English speakers, we paid particular attention to potential ambiguities in meaning and interpretation that could arise due to linguistic nuances. For example, because we noticed that specific perpetrators were mentioned extremely rarely, we decided against the annotation of any abstract entities such as "sexual violence in armed conflict," in order to reflect the low frequency of instances when particular groups or persons are held accountable. Additionally, we adjusted our definition of a "Hero" in order to include those "who ha[ve] the potential to save others, given their abilities, knowledge, positionality/perspective, but who [are] not able to do so because of structural discrimination," in order to account for a common pattern we noticed in the

²We found the names of the debates [here](#).




Figure 1: Character label distribution.

description of women.

Furthermore, similar to [Frermann et al. \(2023\)](#), we allowed for overlapping spans, meaning that entities could take on multiple roles or that an entity appearing within the span of another entity’s definition could still be annotated independently.

After consolidating our annotation guidelines through this initial stage, we again randomly selected 50 additional speeches, aiming to keep a distribution of the representation of years consistent with the make-up of the entire dataset. The total number of files we annotated is therefore 54. Of these, 44 speeches were used as training data, five as development data, and five as testing data for our computational model. Similarly to the first phase in our annotation process, all annotations were first done individually by the authors using [Klie et al. \(2018\)](#)’s annotation platform, INCEpTION.³ The resulting mean inter-annotator agreement across all speeches was 0.38 (Krippendorff’s α), as computed by INCEpTION’s built-in functionality.⁴ After individually annotating all speeches, we curated character roles collectively, deciding on a single set of gold labels.

All curated speeches were then downloaded as .tsv-files, which include the character offsets, the token and its annotation.

Figure 1 represents the frequency of the roles within the annotated data set.

4 Exploratory Work

4.1 Meta-Analysis

Since one of the goals of the WPS agenda is the inclusion of women in peacekeeping processes, we did not only want to analyze textual data but also

³INCEpTION on Github

⁴A detailed overview of pairwise agreement scores between individual annotators and our gold labels can be found in [Appendix B](#).

examine the composition of the Security Council itself. This led to our first research question: "What does the gender representation of the UNSC look like?"




Figure 2: Number of Speeches by Gender Across Different Years (All Speakers)

Using the meta-information extracted from the speech corpora, Table 2 shows that between 2000 and 2019, there was never a time when women were responsible for at least half of all speeches. The gender of some speakers is marked as unknown, primarily in the case of the Secretary-General (who has never been a woman) or the Deputy Secretary-General (predominantly women during this time period), as these roles are typically referred to only by their title rather than with gendered honorifics such as ‘Mr.’ or ‘Ms.’. While the plot shows that overall female participation increased from under 20% in 2000 to almost 40% in 2019—with a peak of just over 40% in 2018—the participation is still fluctuating.




Figure 3: Gender Distribution of Speakers Among Permanent Members of the UNSC

Focusing on speakers from the permanent member countries (China, Russia, USA, UK, and France) in Figure 3, the differences are striking.

While over 80% of the speeches by the US were delivered by female speakers, the other four countries lag far behind. The UK has the second-highest percentage of female speakers at only approximately 20%, while France and Russia are even lower, and China did not have a single female representative. This disparity raises the question of whether women’s representation influences how gender is discussed in the Council. Due to time and space constraints, we will not systematically differentiate between countries in our analysis of narratives, but this would certainly be an interesting topic for future research.

4.2 Preliminary Analyses

As part of a preliminary analysis, we examined the explicit mentions of gendered entities. Specifically, we queried the words "women", "men", "girls", and "boys" in the speeches. Our findings reveal that female entities are mentioned much more frequently than male ones. As Table 1 shows, the number of sentences containing "women" or "woman" is approximately 20 times higher than those containing "men" or "man". A similar trend is observed for children, though to a lesser extent: Sentences containing "girl" or "girls" are about ten times more frequent than those mentioning "boy" or "boys".

Term	Count (%)
"Men" or "man"	1488
"Women" or "woman"	31791
"Men" or "man" (no "woman/women")	313 (21.03%)
"Women" or "woman" (no "man/men")	30616 (96.30%)

Table 1: Occurrences of Gendered Terms Related to Men and Women

Another notable observation is the co-occurrence of gendered terms. When "man" or "men" appear in a sentence, they co-occur with "woman" or "women" in 78.97% of cases. In contrast, "woman" or "women" co-occur with "man" or "men" in only 3.7% of their sentences. A similar pattern is observed for children: when "boy" or "boys" are mentioned, they co-occur with "girl" or "girls" in 65.96% of cases, while "girl" or "girls" co-occur with their male counterpart in just 5.53% of their sentences. This indicates that, although boys are more frequently mentioned alone than men, the difference to the mentioning of girls is still striking.

As expected, and as suggested by the name of the agenda, the primary focus of these debates is on women and girls. This is also reflected in the

Term	Count (%)
"Boys" or "boy"	379
"Girls" or "girl"	4518
"Boy" or "boys" (no "girl/girls")	129 (34.04%)
"Girl" or "girls" (no "boy/boys")	4268 (94.47%)

Table 2: Occurrences of Gendered Terms Related to Boys and Girls

queried terms; since we searched for "women", the results naturally align with the agenda’s title, "Women, Peace, and Security", which likely contributes to the high frequency of female mentions. Although women and girls are mentioned to a much greater extent, men and boys are not excluded from the debate. As a next step, our goal is to move beyond mere frequency analyses and explore the contexts and narratives in which these gendered entities are mentioned.

4.3 Topics

4.3.1 Bertopic

To analyze the contexts in which gendered entities are mentioned, we used BERTopic (Grootendorst, 2022) on the sentences where "women", "woman", "girls", and "girl" occur (without "man", "men", "boy" and "boys") and vice versa. BERTopic (Grootendorst, 2022) utilizes transformers and c-TF-IDF to form dense clusters of texts. These clusters represent topics, which are described by representative keywords to facilitate interpretation.

Overall, we found that BERTopic with stop words identified significantly more topics for female entities (309) than for male entities (7). To simplify interpretation, we used BERTopic’s option to reduce the number of topics of the female entities from 309 to nine. We then created word clouds for the remaining topics for the female and the male entities. Below, we present a selection of interesting topics:

Female entities Among the topics related to female entities, two predominantly revolve around formalities, featuring keywords like "director," "statements," "delivered," "friends," "thank," "Ms.," and "Secretary-General." Another topic is distinctly centered on the WPS agenda, with keywords such as "women," "security," "peace," "resolution," "agenda," "conflict," and "UN-Women." For these three topics, the emphasis is less on specific women and more on the WPS agenda itself. Other topics highlight areas where women are primarily mentioned as victims, such as sexual violence (Figure 4), human trafficking (Figure 5), and

domestic violence (Figure 6). In addition, the Se-



Figure 4: Female Entities: Sexual violence



Figure 5: Female Entities: Human trafficking



Figure 6: Female Entities: Domestic Violence

curity Council frequently references specific terror organizations, such as Chibok and Boko Haram, and the experiences of girls in Nigeria who have been victims of abductions and to other atrocities (Figure 7). Lastly, women are often mentioned in connection with UN peacekeepers or blue helmets (Figure 8). In these instances, women are either part of a narrative advocating for greater female participation in peacekeeping troops, framing them as potential heroes, or linked to another, albeit related, context. This second narrative involves reports of women being victims of sexual violations by (male) peacekeepers, which has led to a loss of trust in peacekeepers and subsequent calls for more female peacekeepers. Thus, women are simultaneously




Figure 7: Female Entities: Terror Organizations

framed as heroes (peacekeepers) and victims (of sexual violations by peacekeepers), while peacekeepers are often framed as villains. Two examples for these framings are:

- (1) "We are convinced of the advantages inherent in increasing the number of women Blue Helmets in peace operations, and we are committed to that cause."
(UNSC_2019_SPV.8649_spch030.txt, Peru)
- (2) "In addition to more women wearing blue helmets, we must take every possible step to root out sexual exploitation and abuse by peacekeepers, which is too prevalent."
(UNSC_2015_SPV.7533_spch015.txt, USA)

It is worth noting, however, that although these cases often imply that the peacekeepers are male, this is not explicitly stated, as can be seen in example (2).



Figure 8: Female Entities: Blue Helmets

Male entities For male entities, one of the topics also involves sexual violence (Figure 9). In these



Figure 9: Male Entities: Sexual violence

sentences, men are framed both as perpetrators and as victims of sexual violence, as seen in the following example:

- (3) "It will take place on 8 or 9 July and will focus on the role of men as perpetrators but also as victims of sexual violence in conflict."
(UNSC_2013_SPV.6984_spch057.txt, Liechtenstein)

conflict zones, perpetrated by men and terrorist organizations like ISIS. Some of these recollections are also present in the topic represented by Figure 12, although the speaker here also talks about specific men they frame as leaders or heroes (and themselves) and also of the importance of including men in gender equality efforts and also highlighting advantages for them as in the following example:

- (4) "Then, in 2009, knowing that the mentality of men must be changed, I established the Network of Men Leaders. I reached out to a great number of men leaders in Governments and the business community, and religious leaders."
(UNSC_2016_SPV.7642_spch019.txt, The Secretary-General)

Thus, this topic is more ambiguous than the others, making it difficult to identify an overarching description for it as it encompasses several topics and framings.




Figure 10: Male Entities: Promotion of Gender Equality

Another topic relates to gender equality (Figure 10), emphasizing the need to include men in efforts to achieve gender equality and highlighting the benefits for them as well.



Figure 11: Male Entities: Victim's reports

The topic represented by the word cloud in Figure 11 primarily includes personal accounts and reports of abuse by (female and male) victims in



Figure 12: Male Entities: Overlapping topics (personal recollections, gender equality, men as leaders)

This analysis revealed that while women are predominantly framed as victims of abuse, trafficking, and violence, they are also frequently portrayed as (potential) heroes in peacekeeping contexts. Men, on the other hand, are mentioned far less frequently and with less thematic variety. Depending on the context, men are framed as victims, heroes, or villains. Notably, these roles are not mutually exclusive; men framed as villains are often contextualized within the circumstances that led to their actions. Furthermore, we observed that when women are mentioned, it is often without male entities, whereas when men are mentioned, it is frequently in co-occurrence with women—for instance, when men are mentioned as perpetrators and women as their victims, but also when both are mentioned as victims together. We have also seen that, although it is often implied that men are perpetrators, as




Figure 13: Three dimensional visualization of the embeddings that we extracted and applied the UMAP dimensionality reduction algorithm to. Each point indicates a speech in our dataset.

in the case of peacekeepers, this is not explicitly stated. This lack of explicit mention can influence analyses that rely on general terms such as "man" or "men."

4.3.2 Extracting Embeddings and Clustering

In addition to topic modeling with BERTopic, we extracted dense semantic embeddings using the mx-bai-embed-large model (Li and Li, 2023), a state-of-the-art embedding model that outputs 1024-dimensional vectors. We obtained document-level embeddings for each full speech in the WPS corpus, allowing us to investigate more general patterns in the discourse.

To visualize and explore the structure of these high-dimensional embeddings, we applied the Uniform Manifold Approximation and Projection (UMAP) dimensionality reduction algorithm (McInnes et al., 2018). We opted for UMAP in virtue of its preservation of both local and global relationships in the data. In this way, we hoped to identify similarity across speeches that could potentially reflect shared rhetorical framing, thematic focus, or policy alignment. We created both two- and three-dimensional visualizations of all of our data points (see Figure 13). In our first renderings of these graphs, we noticed that there were a few highly concentrated points. Upon investigation, we realized that they were the President’s speeches, that contained very little information useful for our purposes, so we removed these from further analyses.

Using the two-dimensional visualization, we applied the Density-Based Spatial Clustering of Applications with Noise (DBSCAN) algorithm (Ester et al., 1996) to discover clusters of semantically similar speeches. We decided to use the DBSCAN algorithm for two key reasons. First, unlike partition-based algorithms such as k -means, DBSCAN does not require the number of clusters to be predefined. This feature is particularly important in the context of exploratory, unsupervised analyses. Second, DBSCAN is particularly well-suited for high-dimensional data because it detects outliers that do not strongly associate with any cluster (Schubert et al., 2017). For each data point, the algorithm constructs an ϵ -neighborhood, which includes all other points within a distance of ϵ from the target point (Géron, 2022). If a point has at least min_samples neighbors within this ϵ -radius, then it is classified as a "core instance." Clusters are thus constructed from continuous sequences of neighboring core instances. If a point is neither a core instance, nor a point in the neighborhood of a core instance, it is classified as an anomaly. These features of DBSCAN allow for extensive exploration identification of clusters that vary in size, while ensuring that outliers are isolated accordingly.

In our exploration of the data, we used standard DBSCAN parameters ($\text{eps} = 0.5$, $\text{min_samples} = 5$), but this configuration failed to yield meaningful clusters due to the sparse distribution of the speech embeddings in our reduced space. After a series of exploratory tuning steps, we increased the neighborhood size to $\text{eps} = 6.5$ while keeping $\text{min_samples} = 5$. This value was selected based on its ability to produce a manageable number of clusters that were both semantically coherent and analytically meaningful. Smaller eps values resulted in excessive fragmentation, with few or no clusters being formed. Conversely, increasing eps beyond 6.5 led to the merging of semantically distinct points into broad clusters that lacked internal differentiation, reducing the algorithm’s ability to distinguish between distinct discursive themes or national patterns. In this context, $\text{eps} = 6.5$ represented a practical upper bound that preserved cluster interpretability while still capturing key structural and semantic differences in the data. With these hyperparameters, DBSCAN identified 18 distinct clusters (see Figure 14). These clusters revealed two overarching patterns:

Country-specific clusters(14/18) The majority of the clusters included speeches from the same




Figure 14: Clusters as identified using DBSCAN. Each point represents a speech. Points that are square shaped with a white outline represent outliers, while circle-shaped points represent speeches that are part of different clusters, distinguished by the color of the outline.

country, suggesting consistent national narratives across multiple sessions. Some examples include Cluster 0 (Chile), Cluster 2 (Iceland), and Cluster 11 (Italy). Since text embeddings are most sensitive to semantic patterns, it is unsurprising that several of these clusters reflected the use of recurring rhetorical templates. In Cluster 11, for instance, three of the seven speeches began with the following identical phrase: *"Allow me to begin by thanking you, Mr. President, for organizing this open debate of the Security Council."* These types of patterns underscore the highly formulaic nature of diplomatic discourse in international forums such as the UN, where delegations often adhere to established conventions in structure and content. Therefore, the emergence of these clusters seems to reflect both the semantic coherence of the messaging provided by different countries, as well as the institutionalized norms in public speaking in multilateral settings. Beyond these standard diplomatic courtesies, however, we also notice that entire phrases such as *"Women are formidable negotiators, mediators, and peacebuilders"* and *"Sexual violence is not cost-free"* are repeated word for word across multiple statements. These overlaps suggest that some speeches may have been prepared from standardized templates or adapted with only minor edits over time, reflecting the institutional inertia of multilateral diplomacy. Overall, these examples not only point to the formalized norms of speech delivery in international forums like the UN but also raise questions about the performative versus substantive nature of these public commitments.

Thematic clusters (4/18) A smaller but analytically important subset of clusters grouped together speeches that focused on similar substantive issues, regardless of the speaker's national affiliation, year, or session. This indicates thematic convergence of all countries on items that are highly salient on the WPS agenda. Here, a clear example is cluster 10, which includes speeches from a diverse set of countries: Gabon, Mexico, Australia, South Africa, Lithuania, Namibia, Rwanda, Morocco, the Russian Federation, the Republic of Korea, Ghana, and Egypt. Despite the differences in their geopolitical stances, representatives from these countries all delivered speeches sharing the thematic focus of sexual violence in armed conflict, a highly salient topic within WPS debates. Visual inspection of the clusters confirms that the phrase *"sexual violence in conflict"* or *"sexual violence in armed conflict"* appears repeatedly across the texts. For the most part, the lexical similarity is not superficial; these references are embedded in statements that express gratitude to the UN Secretary-General, commend the work of the Special Representative on Sexual Violence in Conflict, and emphasize the international community's responsibility to act. Several speeches reference the efforts of Special Representatives such as Zainab Hawa Bangura and Pramila Patten, and acknowledge the institutional frameworks set up to address sexual violence as a tactic of war. All speeches converge around the framing of sexual violence as a war crime, a crime against humanity, and a threat to peace and security, and emphasize the continued need to take action. Nevertheless, the inclusion of the Russian Federation in this cluster illustrates that thematic convergence does not necessarily imply full discursive or political alignment. While addressing the same issue, the Russian delegation explicitly voiced concern over the expansion of UN mandate language beyond existing Security Council resolutions, highlighting that even within thematically coherent clusters, substantive differences in framing and emphasis persist.

In addition to these categories of clusters detected by DBSCAN, a substantial number of speeches were marked as outliers. Among these outliers, upon inspection of the two-dimensional UMAP plot, we identified a semantically coherent grouping. Although not dense enough to form a cluster given our DBSCAN hyperparameters, these neighboring points all addressed the problem of human trafficking.

This distribution, where country-specific clusters are more prevalent, but where a notable minority also reflects content-driven convergence, demonstrates how unsupervised methods like DBSCAN can surface both stylistic consistency and shared rhetorical priorities. On the one hand, this approach offers a scalable way to detect patterns in country discourse, revealing when commitments may stem from genuinely sustained national positions, or, conversely, from the mechanical reuse of established templates. On the other hand, such clustering methods are limited in their ability to capture how countries frame issues like sexual violence. They also fail to distinguish between more or less assertive stances, or between expressions of solidarity versus expressions of concern over personal interests.

This limitation highlights the need for more fine-grained analyses, especially when exploring how speakers frame themselves or others as heroes, villains, or victims in these debates. To that end, we turn next to the annotation of character roles within the speeches, building on these clustering and topic modeling results to examine the underlying narrative structure of WPS discourse.

The main objective of this study is to examine how heroes, victims, and villains are portrayed in discourse and to assess whether LLMs can accurately identify these roles. During the annotation process, we observed that detecting such framings can be particularly complex, especially in text genres designed to appear neutral and devoid of emotion. This led us to question whether LLMs might capture subtle cues and implicit signals that human annotators might overlook, and vice versa.

As presented in related works, prior research on the detection of semantic roles such as heroes, victims, and villains has often relied on a combination of NER and sentiment analysis (Gomez-Zara et al., 2018b). More recent studies have leveraged transformer-based architectures such as BERT (Sharma et al., 2023) and GPT (Stammach et al., 2022) for similar tasks. Building on these insights, we conducted our own experiments to explore the effectiveness of different approaches in identifying these roles.

To this end, we employed three distinct methods: (1) a fine-tuned RoBERTa model, and (2) direct prompting of LLMs.




Figure 15: Top Entities per Label

4.4 Results from Annotated Data

Before moving on to the models, we will briefly present some insights from the annotated speeches. These statistics were only generated after fine-tuning RoBERTa and employing the LLMs to prevent any manipulation.

Figure 15 presents the top entities. As already theorized, the most frequently mentioned entity under the "victim" label is "women". However, the top ten also include groups that encompass women, such as "women and girls", "women's", and "women and children". In contrast, the most frequent entity for the "hero" label is "we", referring to the speaker, their countries, and/or the Security Council itself. This aligns with previous findings that member states tend to highlight their own contributions (True and Wiener, 2019).

Additionally, six more references appear to countries, the Security Council, and the Secretary-General, as well as to "Estonia". Interestingly, the second most frequently mentioned hero entity is also "women", suggesting that the debates do not solely frame them as victims but also acknowledge their contributions. This dual portrayal is further evidenced by the fact that "women" is the top entity in the dual-label category of hero and victim, followed by its possessive form "women's" and specific mentions such as "Nadia", "these women", and "women in armed conflict".

Turning to the "villain" category, we observe




Figure 16: Character Labels for "women" and "women's" Across the Years

that this label is assigned far less frequently than the other labels. The most common villain entity is "perpetrators", followed by variations such as "the perpetrators" and "perpetrators of sexual violence" (ranked 7th). This suggests that perpetrators are often left unnamed or vaguely referenced. Another notable observation is that speakers at times seem to villainize themselves, particularly when discussing the inaction of the Security Council:

"Like those who preceded us, we welcome the decision by the presidency to hold this open debate to mark the resolution's third anniversary. Such debates remain necessary, considering how far **we** are from full implementation."
(UNSC_2003_SPV.4852Resumption1_spch035)

Aside from this, only two specific groups are villainized: United Nations peacekeepers and Daesh (ISIS).

The hero-villain dual label is rare but includes references such as "the Security Council" and "a peacekeeper", indicating that while peacekeepers are sent with good intentions, they are sometimes also perpetrators of (sexual) violence, two narratives that we have already seen in BERTopic in the previous section.

Examining entities that contain "women" and "women's" (see Figure 16), we find that although portrayals of women shift over the years, there is no clear trend indicating a transition from victimhood to agency, as observed in [de la Rosa and Lázaro \(2019\)](#) for UN resolutions.

Further analysis (see [Appendix A](#)) reveals that "women" alone constitutes the largest group within hero, victim, and hero-victim labels. In the victim category, one-fourth of occurrences involve "women" in conjunction with "girls" or "children",




Figure 17: Character Labels for "men" and "men's" Across the Years

reinforcing findings from ([Carpenter, 2005](#)) that these groups are often used synonymously with civilians in general. Meanwhile, villain mentions of "women" mostly appear as part of longer noun phrases describing perpetrators, as seen in this example:

- (5) **"Those who commit crimes against women, including the peacekeeping personnel, should be brought to book."**
(UNSC_2000_SPV.4208Resumption2_spch004.txt)

This also illustrates how perpetrators are vaguely referenced—whether deliberately or due to a lack of information. There are never any female perpetrators mentioned.

Turning to the character labels for "men" and "men's" (Figure 17), we again see that men are mentioned far less frequently than women across years. Interestingly, men appear in multiple roles (hero, hero-victim, victim, hero-villain) in 2000, 2002, and 2013, but only as heroes in 2003 and 2004, and only as victims in 2012, 2014, 2016, and 2018. This could suggest increasing recognition of male victims of armed conflict, as well as a reduction in the portrayal of men as perpetrators.

A pie chart in [Appendix A](#) reveals that among entities containing "men" or "men's", 42.9% of hero-labeled mentions refer solely to "men", while the remainder also include "women". For victim-labeled mentions, only 22% reference men without also mentioning "women" or "boys", indicating that when men are framed as victims, they are rarely the sole victims but rather appear alongside women or children. This supports our broader observation that men are less frequently discussed than women and often in conjunction with them.

These insights from our labeled entities highlight key framing patterns in the speeches: women

are most frequently mentioned as victims, but they also appear as heroes; villains tend to be vaguely referenced; and men, when mentioned, are often grouped with women or children. With this contextual understanding, we now turn to our models to assess how well they can classify and analyze these patterns.

4.5 RoBERTa Model

We chose to fine-tune a BERT-based model as part of our approach, drawing inspiration from prior work on similar tasks. Various BERT models were employed by participants in a shared task aimed at identifying whether characters in memes were portrayed as glorified, vilified, or victimized (Sharma et al., 2022). Notably, both the top-performing system (Kun et al., 2022) and the third-ranked system (Singh et al., 2022) incorporated RoBERTa (Liu et al., 2019) into their approaches. Building on these insights, we opted to fine-tune RoBERTa for our task. While exploring alternative architectures could have provided additional insights, our primary objective was to establish an initial approach and evaluate whether an existing, well-performing model could effectively capture the roles of heroes, victims, and villains in the WPS. Future work could extend this by experimenting with a broader range of model architectures.

4.5.1 Preprocessing steps

To adapt the NER task to a transformer-based model, where tokens with labels do not stand alone but rather belong to sequences, we applied the BIO labeling scheme (Tjong Kim Sang and Buchholz, 2000). The BIO scheme assigns labels as follows: the first token of an entity sequence is prefixed with "B-", subsequent tokens in the sequence receive an "I-" prefix, and tokens without any annotation are labeled "O". Table 3 illustrates an example of this labeling scheme applied to a sample sentence.

For tokens, where entities overlapped, we modified the existing separate labels into a combined new label in order to avoid a multi-label classification task. We decided to proceed like this since the occurrence of overlapping labels is rather underrepresented (see Figure 1) and would not lead to promising results. This modification results in a total amount of 13 labels⁵.

⁵Labels: HERO, VICTIM, VILLAIN, HERO_VICTIM, HERO_VILLAIN, VICTIM_VILLAIN with B- and I- preceding respectively, as well as O.

4.5.2 Implementation

To fine-tune the RoBERTa model for the multi-class classification, we followed standard implementation steps, including encoding labels and pre-tokenised input text using pre-trained RoBERTa embeddings. For batch training, we set a batch size of 16, ensuring a balance between computational efficiency and model performance.

To determine suitable hyperparameters, RoBERTa was fine-tuned on the training dataset using six different hyperparameter configurations (see Appendix D) and evaluated on the development dataset. As the primary objective of this implementation is to assess whether the model can identify roles given the constraints of a small dataset, the choice of hyperparameters was deliberately limited to ensure computational feasibility within the available resources. The varied hyperparameters included the learning rate (1e-5, 3e-5 or 5e-5), the number of training epochs (3 or 30) and the choice of optimizer (AdamW or Stochastic Gradient Descent). The performance of the fine-tuned models on the development dataset is shown in Table 4, based on their F1 values for models m1-m4. Models m5 and m6 are excluded from the table as their results were identical to those of m4. The following chapter about evaluation (4.5.3) will explain how exactly the evaluation was conducted.

	F1				S
	m1	m2	m3	m4	
HERO	0.59	0.62	0.63	0.64	122
VICTIM	0.67	0.72	0.73	0.73	77
VILLAIN	0.54	0.52	0.5	0.52	10
HERO_VICTIM	0.45	0.46	0.36	0.4	26
HERO_VILLAIN	0	0	0	0	0
VICTIM_VILLAIN	0	0	0	0	3
micro avg	0.6	0.62	0.62	0.64	238
macro avg	0.45	0.46	0.44	0.38	238
weighted avg	0.59	0.62	0.62	0.63	238

Table 4: F1 scores of the models m1, m2, m3, and m4. The models m5 and m6 are not included because they result with the same evaluation scores as m4. S represents the number of instances per class in the dataset.

The results presented in Table 4 indicate that models m3 and m4 achieve the highest F1 scores for the two most frequently occurring classes, heroes (0.63 for m3 and 0.64 for m4) and victims (0.73 for both m3 and m4). Furthermore, model m4 yields the highest weighted average F1 score (0.63), which accounts for the support of each class when

Token	"[...]"	my	country	has	[...]	defending	the	victims	of	terrorism	."
Labels	O	HERO	HERO	O	O	O	VICTIM	VICTIM	VICTIM	VICTIM	O
BIO Labels	O	B-HERO	I-HERO	O	O	O	B-VICTIM	I-VICTIM	I-VICTIM	I-VICTIM	O

Table 3: Example of the BIO annotation scheme applied to a sample sentence from UNSC_2015_SPV.7585_spch010.txt. The whole sentence is: *Council members know that my country has been a standard-bearing country for defending the victims of terrorism.*

computing the overall performance. These findings served as the primary rationale for selecting the hyperparameters of m4 for the final fine-tuning process.

Accordingly, the final model was fine-tuned using both the training and development datasets.

4.5.3 Evaluation & Interpretation

To evaluate the prediction results of RoBERTa — both during hyperparameter optimization and final testing — we used *sequeval*⁶ by Nakayama (2018), a framework designed for the evaluation of sequence labeling tasks. We employed the default evaluation mode, which allows for some leniency in labeling. In contrast, *sequeval* also offers a strict mode, in which token-level predictions must match exactly; otherwise, an entire sequence prediction is considered incorrect, even if the majority of tokens or all but one are correctly identified.

The performance of the RoBERTa-based classifier on the WPS corpus is summarized in Table 5. The evaluation metrics include precision (P), recall (R), and F1 for each class, along with the S.

The model performs best in detecting the roles of heroes and victims, achieving F1-scores of 0.72 and 0.73, respectively. This indicates a strong ability to recognize these roles. For villains, it has a lower precision (0.43) and F1-score (0.55), despite a relatively high recall (0.75). This suggests that, while the model is able to identify villains when they are present, it also frequently misclassifies other roles as villains, leading to a higher false positive rate. However, the support indicates that there are only four occurrences of the label, leaving the open question if the model’s results would be higher if there was a higher frequency.

Unfortunately, the results on the test dataset do not give a lot insight into its efficiency regarding overlapping labels. The labels HERO_VILLAIN and VICTIM_VILLAIN are not represented in the test data. But we do get some insight due to the represented label HERO_VICTIM with 12 occurrences. The model struggles with the

	P	R	F1	S
HERO	0.69	0.76	0.72	95
VICTIM	0.67	0.81	0.73	32
VILLAIN	0.43	0.75	0.55	4
HERO_VICTIM	0.36	0.33	0.35	12
micro avg	0.65	0.73	0.69	143
macro avg	0.54	0.66	0.59	143
weighted avg	0.65	0.73	0.69	143

Table 5: Evaluation results of the fine-tuned RoBERTa model. Precision (P), recall (R) and F1 results for each label, for both micro and macro average, as well as their support (S).

HERO_VICTIM category, achieving the lowest F1-score (0.35). This suggests difficulty in distinguishing overlapping roles, likely due to ambiguity in textual representations or annotation challenges.

The micro and macro averages give insight into the overall performance. The micro-averaged F1-score of 0.69 reflects a relatively strong overall performance, while the macro-averaged F1-score of 0.59 indicates class imbalance.

The weighted average F1 score of 0.69 is close to the micro-average, confirming that the model performs well overall when considering the class distribution. This suggests that the model’s strengths in detecting the more common categories (HERO and VICTIM) help to maintain a solid overall performance. However, as the weighted average gives more weight to classes with higher support, it does not fully reflect the difficulties the model has with under-represented labels such as VILLAIN and HERO_VICTIM.

The discrepancy between the macro (0.59) and weighted (0.69) F1 scores highlights the impact of class imbalance. While the model performs well with the dominant classes, its struggles with less common labels suggest a potential bias towards the majority classes. Addressing this imbalance - perhaps by augmenting the data, reweighting the loss functions, or using different sampling techniques - could help to improve performance on the rarer categories and produce a more balanced model.

⁶<https://github.com/chakki-works/sequeval>

Overall, this model performed exceptionally well, with its predictions closely aligning with our annotations in most cases. The consistency of its outputs suggests a high degree of reliability in capturing the intended role assignments. Interestingly, the models also identified some entities that we annotators had overlooked, indicating their potential to enhance annotation quality and reduce human error. This suggests that these models could serve as valuable tools for automating or assisting in large-scale annotation efforts, provided that their occasional inconsistencies are accounted for through further refinement or human verification.

However, certain inconsistencies were observed, such as instances where a model initially assigned one label but later switched to another or where only a partial entity, such as "the", was annotated instead of the full noun phrase.

A further error analysis will follow.

4.6 LLMs

Following a similar methodology as [Stammbach et al. \(2022\)](#), we prompted to two LLMs and evaluated their outputs. We analyzed the results of DeepSeek-R1 (DeepSeek) ([Guo et al., 2025](#)) and Meta-Llama-3-70B (Llama) with both **zero-shot** and **few-shot** prompting.

DeepSeek is a large language model specifically designed for reasoning, making it particularly relevant for our study, as one of our main objectives was to analyze how it draws conclusions. The model is enhanced through pure reinforcement learning, built on the fine-tuned DeepSeek-V3-Base model, and utilized with a reinforcement learning framework. Notably, it does not rely on supervised fine-tuning, which enables chain-of-thought reasoning and enhances its overall inferential capabilities. Initially, we intended to compare DeepSeek with OpenAI’s o1-1217 model, as both have demonstrated comparable performance ([Guo et al., 2025](#)). However, due to access limitations, we ultimately opted to compare DeepSeek with Llama, a more accessible alternative to OpenAI’s models.

Llama is a state-of-the-art open-weight transformer-based model developed by Meta, designed to perform a wide range of natural language processing tasks. It has demonstrated strong generalization capabilities across various benchmarks, making it a suitable candidate for the comparative analysis ([Grattafiori et al., 2024](#)).

4.6.1 Preprocessing

In order to ensure that our models would properly deal with the long speeches and guidelines that we provided in the prompt, we opted for requesting that models reproduce the input speech, solely adding in-line beginning and end tags (e. g., `<HER>Ambassador Powers</HER>` represents the character annotation for a Hero). Given this initial format, we converted both our own annotations and the model outputs into span-based JSON files.

Due to some formatting issues with whitespace and newline characters and span assignment in our gold annotations, we developed a process of matching characters and checking if their initial offsets (start and end markers) as extracted from the LLM output corresponds to our own. We first extracted the entities included inside the beginning and end tags, maintaining their offsets from the LLM output. We used `difflib.SequenceMatcher`⁷ to create a fuzzy matching algorithm, setting a similarity threshold of 0.98. In parallel, for each speech, we constructed a map between the character offsets in our initial texts and those returned by the LLMs. If any discrepancy in characters corresponding to the same offset was detected, we adjusted the map by either incrementing or decreasing the offset accordingly. We then used this offset map to recalculate the span boundaries for all of the extracted entities, adding spaces where necessary in order to match with the span in the original texts. This enabled us to accurately compare tagged entities according to exact offsets in the text.

4.6.2 Implementation

Since OpenAI’s models were not available for testing, we instead used ChatGPT for prompt engineering. We started with a prompt, evaluated the output based on the speeches in the dev set, and continuously modified the prompt to align with the constructions. The resulting prompts include the task description followed by the annotation guidelines ([Appendix C](#)), a defined format of what the output should look like, and a short repetition of the task to support long-text understanding. While these are all the information the zero-shot prompt gets, the few-shot prompt also includes two annotated speeches from the train data as examples. The **zero-shot** prompt is available in [Appendix E](#) and the **few-shot** prompt in [Appendix F](#).

⁷<https://github.com/python/cpython/blob/3.13/Lib/difflib.py>

Since DeepSeek is a reasoning model, we considered that it generates longer outputs than Llama, hence the number of output tokens as response from the LLMs was maximized.

4.6.3 Evaluation & Interpretation

We present our results in Table 6. Although we did not use the same train-development-testing split as we did for RoBERTa, we still only tested the models' performance on the speeches we deemed as part of the testing set, in order to facilitate comparisons of both quantitative and qualitative measures. Our evaluation framework included both exact and partial match criteria. In line with our reporting of more lenient evaluation results from RoBERTa, we also only report the results of our partial matching method. While exact matches required full alignment in both character span and role label, we define partial matches as having at least a 50% overlap in start and end spans between LLM-generated outputs and our gold labels. Additionally, we added the constraint that a gold-label span could not be matched to multiple predicted spans, in order to avoid inflated performance scores. This means that even with our partial matching strategy, models are still penalized if they split lists into separate individual characters, which was one of the most prominent guidelines that we provided models with.

A key distinction between the evaluation framework we developed for LLMs and the one we developed for RoBERTa concerns the treatment of multi-label annotations for the same span. For the implementation of RoBERTa, we conceptualized this task as a single-label classification task, introducing composite class labels, and asking the model to provide predictions that represent all co-occurring roles. This approach simplifies evaluation, but, as we noticed, it also introduces label sparsity and reduces generalizability. In contrast, in our LLM set up, we preserve each role label as an independent annotation, even when multiple labels are assigned to the same exact span, or when two distinct labels have overlapping spans. We treat each character as distinguished by its start, end, and gold label, as a distinct evaluation unit. Thus, "credit" is assigned per label, and not conditioned on jointly predicting all labels for a given span.

While we acknowledge that a strict comparison between the two LLMs that we utilized is not entirely fair in virtue of the immense difference in parameters, in this section, we would still like to

comparatively report differences in their performance.

As seen in Table 6, across all labels, DeepSeek consistently outperforms Llama in terms of micro-averaged and weighted scores, especially in the few-shot condition, suggesting that it is better able to generalize from limited examples.

For the Hero label, DeepSeek in the zeroshot setting demonstrates the strongest recall across all characters and model configurations (0.760). Nevertheless, this negatively affected its precision score, which, at a low 0.392, negatively impacted the overall F1 score (0.517). Few-shot prompting improved both precision (0.478) and F1 (0.575). Llama, however, performs worse, with a 0.464 score for this label in the few-shot condition, and an even worse score for the zero-shot condition. There is a similar pattern for the Victim label, where DeepSeek with few-shot prompting achieves the highest F1 score (0.705), driven by a high precision of 0.738.

Performance for the prediction of the Villain tag, however, reveals an unexpected pattern. Llama with zero-shot prompting achieves the highest scores across the board (0.75 precision, recall, and F1). While these results might seem impressive, it is important to note that our test set only included a total of four Villain labels, and the differences in scores across conditions are in fact driven by differences in the annotation of one or two entities. In future work, we wish to investigate whether this trend in the level of "understanding" of Villains remains when the test set is expanded.

Aggregate performance metrics confirm the better performance of DeepSeek across the board. In terms of micro-averaged F1, which weights performance by the number of true examples, DeepSeek scores 0.607 in the few-shot setting compared to Llama's 0.481. Macro-averaged F1, which treats all classes equally, shows a similar pattern (0.617 vs. 0.518). Most notably, weighted averages, which factor in class imbalance, demonstrate DeepSeek's clear advantage, with a weighted F1 of 0.620 in the few-shot condition compared to 0.498 for Llama. This suggests that DeepSeek not only achieves higher per-label accuracy but also performs better in proportion to the frequency of roles in the dataset.

Overall, we find that the four instances of the models that we used for character annotations perform quite poorly when compared to our expectations and to RoBERTa. In the following section, we

Label	Model	Setting	P	R	F1
Hero	DeepSeek	Zero-shot	0.392	0.760	0.517
		Few-shot	0.478	0.721	0.575
	Llama	Zero-shot	0.292	0.587	0.390
		Few-shot	0.360	0.654	0.464
Victim	DeepSeek	Zero-shot	0.488	0.645	0.556
		Few-shot	0.738	0.674	0.705
	Llama	Zero-shot	0.400	0.652	0.496
		Few-shot	0.429	0.652	0.517
Villain	DeepSeek	Zero-shot	0.333	0.250	0.286
		Few-shot	0.667	0.500	0.571
	Llama	Zero-shot	0.750	0.750	0.750
		Few-shot	0.667	0.500	0.571
Micro avg	DeepSeek	Zero-shot	0.418	0.694	0.522
		Few-shot	0.535	0.701	0.607
	Llama	Zero-shot	0.326	0.610	0.425
		Few-shot	0.382	0.649	0.481
Macro avg	DeepSeek	Zero-shot	0.404	0.552	0.453
		Few-shot	0.627	0.632	0.617
	Llama	Zero-shot	0.481	0.663	0.545
		Few-shot	0.485	0.602	0.518
Weighted avg	DeepSeek	Zero-shot	0.436	0.686	0.533
		Few-shot	0.584	0.662	0.620
	Llama	Zero-shot	0.362	0.627	0.462
		Few-shot	0.418	0.607	0.498

Table 6: Evaluation results (partial match setting) for character-role classification across LLMs. Precision (P), recall (R), and F1 are shown for each label and setting.

present a more detailed description of the mistakes made by each model, offering supporting examples for each common mistake we observed.

4.7 Error Analysis

In the following, we reflect on the strengths, but also the mistakes of the LLMs RoBERTa, DeepSeekZero (DeepSeek with zero-shot prompting), DeepSeekFew (DeepSeek with few-shot prompting), LlamaZero (Llama with zero-shot prompting), and LlamaFew (Llama with few-shot prompting).

4.7.1 Non-Compliance with Annotation Rules

NPs, Articles and Pronouns The annotation guidelines explicitly state in the character-level rules that whole noun phrases (NPs), including restrictive relative clauses, possessive modifiers, numerical descriptors and other descriptive modifiers, must be included when assigning a label to an entity (Appendix C). While this rule is relatively straightforward for human annotators, the performance of language models in complying with this

requirement varies considerably.

RoBERTa is generally successful in capturing tokens that belong to the required NP sequences. However, chat-optimized language models have considerable difficulty in correctly capturing these elements, or in maintaining multiple characters listed together as a single entity when they share a common role or action. A particularly striking problem is the frequent mislabelling, i.e. separation, of the phrase "women and children", which occurs repeatedly throughout the dataset.

Of the models evaluated, LlamaZero and DeepSeekZero most frequently fail to group entities correctly, often omitting articles, pronouns and descriptive modifiers in the labelling process. While the few-shot models show a slight improvement in integrating articles, they still struggle to consistently include relevant pronouns and descriptive modifiers. They also fail to maintain coherent multi-word entity spans.

To illustrate these challenges, the following examples highlight tokens that were incorrectly omit-

ted from the model predictions (marked in bold):

- (6) "**thousands of** women and girls"
(UNSC_2011_SPV.6642_spch055.txt)
- (7) "**the** security council"
(UNSC_2009_SPV.6195_spch008.txt)
- (8) "women **in armed conflict**"
(UNSC_2009_SPV.6195_spch008.txt)
- (9) "**the office of the personal representative of the head of state**"
(UNSC_2016_SPV.7704_spch087.txt)

Conversely, in some cases the models showed overly greedy behavior, incorrectly including tokens that were not part of the intended entity space. These cases are also highlighted in bold below:

- (10) "*the Government's* **own resources**"
(UNSC_2016_SPV.7704_spch087.txt)
- (11) "**the zero tolerance policy of the secretary-general**"
(UNSC_2011_SPV.6642_spch055.txt)

Additionally, several cases were observed where entities were mistakenly separated instead of being recognized as a single cohesive unit. Some examples include:

- (12) "government" and "United Nations system" instead of "*the government and the United Nations system*"
(UNSC_2009_SPV.6195_spch008.txt)
- (13) "my country" and "DRK" instead of "*my country, the Democratic Republic of Congo*"
- (14) "survivors" and "demobilized fighters" instead of "*survivors and demobilized fighters*"
(UNSC_2016_SPV.7704_spch087.txt)

Overall, it remains unclear why the prompting-based models exhibit such discrepancies. However, one possible explanation is that the prompts provided to these models may be too long, as some models struggle with long-text understanding (Wu et al., 2024). This hypothesis is further supported by the observation that few-shot prompting, which includes additional examples toward the end of the prompt reinforcing the annotation rule, tends to yield better results than zero-shot prompting.

Handling Spelling Errors In our test dataset, we identified a spelling error in the following sentence:

- (15) "In this regard, in 2015, 111 cases were decided by the military justice system without regard to the rank **ofthe accused.**"
UNSC_2016_SPV.7704_spch087.txt

During the curation process, we identified *accused* as a villain. As we aim to annotate entire noun phrases (NPs), we would need to include the article *the* as well. But due to the spelling mistake in the original data, the tokens *of* and *the* were not separated, which confronted us with the question of how to label this entity. Ultimately, we decided to annotate *ofthe accused* as a whole entity with the intention of analyzing how different models handle such cases.

Looking at the results generated by the LLMs, we found that none of the models, except RoBERTa, were able to assign a label to the entity at all. DeepSeekZero, DeepSeekFew, LlamaZero and LlamaFew failed to recognize the entity at all, while RoBERTa correctly identified the entity per se. However, it misclassified it as the victim rather than the villain, although it did at least recognize the entity. The question remains as to why *ofthe* was not annotated. As shown in section 4.7.1, RoBERTa does not generally struggle to correctly annotate articles, pronouns or whole noun phrases. Therefore, we hypothesize that the model ignored *ofthe* due to its misspelling.

Welcomed People and Speakers A relatively fixed component of WPS agenda speeches is the acknowledgement of other council members through welcoming remarks and expressions of gratitude. While our annotation guidelines primarily label individuals being thanked as heroes, the LLMs showed inconsistent labelling behaviour.

RoBERTa occasionally labelled people who were greeted as heroes. DeepSeekZero, however, went beyond this by annotating meta-information such as the name of the speaker at the beginning of a document. Furthermore, despite explicit instructions in the annotation guidelines that entities expressing gratitude should not be considered heroic, both DeepSeekZero and DeepSeekFew incorrectly labelled such entities as heroes.

Surprisingly, the results from Llama do not do such a thing and correctly do not label thanking and welcoming and welcomed people.

Resolutions Resolutions, along with other forms of legal instruments, are frequently referenced in WPS speeches. In some cases, they are personified and depicted as active agents, while in others, they are mentioned passively without any agency. To maintain consistency in our annotations, we established a rule in the guidelines specifying that only resolutions explicitly portrayed as actively performing an action should be annotated, while all other types of legal documents should be excluded.

RoBERTa and LlamaZero generally adhered to this rule, correctly distinguishing between active and passive references. However, the results from DeepSeek models and LlamaFew deviated from this guideline, as they frequently annotated resolutions as heroes even when they were merely mentioned as having been adopted⁸ or described as worthy of implementation⁹. While one might argue that resolutions could be considered heroes if they function as active entities, LlamaFew further misapplied the annotation by labeling "the Dakar Declaration" as a hero¹⁰, disregarding the guideline that only resolutions—not other legal documents—should be annotated.

Multi-Labels To further assess the performance of the models, we examined their handling of multi-label annotations, where entities can have multiple roles simultaneously. The results varied between models, with some showing partial success in capturing multi-label assignments, while others introduced errors or inconsistencies.

RoBERTa was able to predict multi-label annotations in certain instances but also made incorrect predictions. For example, it occasionally assigned multiple roles correctly but struggled with consistency¹¹. This suggests that while RoBERTa recognizes multi-label entities to some extent, it lacks a systematic approach to applying them accurately.

DeepSeekZero performed well in handling multi-label annotations in certain cases. Sometimes it correctly assigned both victim and hero labels to women¹². This demonstrates that, under certain conditions, the model can effectively capture the complexity of role attributions, especially in contexts where women are portrayed as both empowered actors and affected individuals.

⁸UNSC_2009_SPV.6195_spch008.txt

⁹UNSC_2011_SPV.6642_spch055.txt

¹⁰UNSC_2011_SPV.6642_spch055.txt

¹¹UNSC_2004_SPV.5066Resumption1_spch026 and
UNSC_2011_SPV.6642_spch055.txt

¹²UNSC_2013_SPV.7044_spch013.txt

LlamaFew showed mixed results. It correctly assigned both hero and victim roles to *women's rights* and *women's empowerment*, in line with the nuanced framing of these concepts¹³. However, it misclassified *strengthening the role of women in conflict resolution and prevention* as both hero and victim, whereas our guidelines dictate that this phrase should only be annotated as hero. This suggests that while LlamaFew has a stronger ability to recognize multi-label scenarios, it may also over-generalize, leading to role assignments that deviate from the annotation guidelines.

Overall, these results highlight the varying effectiveness of different models in multi-label classification. While some models, such as DeepSeekZero and LlamaFew, demonstrated an ability to handle multiple roles for an entity, inconsistencies in application suggest that further refinements are needed to improve the robustness of multi-label role assignment.

Bias During our analysis of the model predictions, we identified several potential biases in the outputs of RoBERTa and LlamaZero, which may be due to underlying patterns in the models' training data or fine-tuning processes.

RoBERTa showed a possible bias in its treatment of the term *council*. The model frequently annotated *council* as a hero, likely due to fine-tuning data in which it appeared predominantly in the context of *council members*. This bias was particularly evident when the model later annotated *the Security Council* as a hero in a context where the more appropriate annotation would have been the adoption of a resolution. Similarly, RoBERTa assigned the hero role to *the Secretary-General* even though the entity was only mentioned in the context of a report, raising concerns about a possible bias in favour of senior officials¹⁴. A similar pattern was seen in the annotation of *my government* following a mention of an official appointment. In addition, RoBERTa ascribed the heroic role to *the unit commanders of the Forces Armées de la République Démocratique du Congo*, even though the sentence merely described their responsibilities and did not portray them as actively solving a problem.

Another worrying trend in RoBERTa's predictions was its tendency to systematically assign the role of victim to women without clear justification. In several contexts where women were mentioned,

¹³UNSC_2011_SPV.6642_spch055.txt

¹⁴UNSC_2016_SPV.7704_spch087.txt

the model labelled them as victims, even though there was no textual evidence to support their victimisation. This suggests a potential gender bias in RoBERTa’s role assignment, reinforcing pre-existing stereotypes.

LlamaZero demonstrated a related but different problem. The model incorrectly annotated the phrase *training must target not only girls and women, but also boys and men*¹⁵ by assigning hero and victim roles to girls and women, while labelling boys and men as heroes only. This inconsistent annotation suggests a gender bias in role attribution, with women being disproportionately assigned the role of victim, while men were primarily portrayed as heroes.

These findings highlight the need for careful consideration of bias when applying LLMs to role classification tasks, particularly in contexts involving gender and institutional entities. Further research is needed to determine whether these biases arise from pre-training data, fine-tuning procedures, or specific annotation guidelines used in previous studies.

4.7.2 Specific Examples of Misclassifications

Sometimes, the span and detection of an entity is correct and a model simply misclassifies it.

- (16) "Within that framework, **local women’s entities** are a driving force in empowering women and strengthening good governance."(*UNSC_2013_SPV.7044_spch013.txt*)

In this example, DeepSeekZero annotated "local women’s entities" as victims and heroes, whereas all other LLMs correctly labeled them as heroes only. The reasoning provided by DeepSeekZero sheds light on this misclassification: *"For example, 'local women’s entities' are a driving force in empowering women, so Hero. But women might also be Victims needing empowerment."* This explanation suggests that the model struggled to differentiate between the first occurrence of "women", where they are described as agents of empowerment, and the second occurrence, where "women" are framed as needing empowerment. Additionally, the model may exhibit a bias toward annotating women as victims, given the frequency of this framing in the speeches. However, since this error does not appear in DeepSeekFew, it highlights that annotations can sometimes be inconsistent or even arbitrary.

Other models also make classification mistakes:

¹⁵UNSC_2004_SPV.5066Resumption1_spch026

- (17) "**Indigenous and local women** have a strong advocacy role, and it should be utilized in all aspects of security and peace."
(*UNSC_2004_SPV.5066Resumption1_spch026.txt*)

In this case, both LlamaZero and LlamaFew misclassified the entity, whereas both DeepSeek models correctly labeled "Indigenous and local women" as heroes. LlamaFew assigned both victim and hero labels to the entity, while LlamaZero performed even worse by incorrectly splitting it into two separate entities—"Indigenous" and "local women"—and assigning both labels to each. This is particularly problematic because "Indigenous" alone is not even a valid entity.

The two annotated speech examples available for LlamaFew may have helped it avoid the same mistake as LlamaZero, but both models still exhibit a bias toward labeling female entities as victims. In this case, interpreting them as victims appears overly speculative—perhaps based on the assumption that their advocacy role has not been sufficiently recognized. This demonstrates how certain biases can lead to overly broad or inconsistent annotations. A similar error can be observed in the following example:

- (18) "My delegation has a special interest in the topic, because resolution 1325 (2000) represented an enormous step forward in the protection of **women** and highlighted the importance of their role in all aspects of United Nations peacekeeping."
(*UNSC_2011_SPV.6642_spch055.txt*)

All models annotated women only as victims, rather than as both victims and heroes. However, the speaker clearly emphasizes their important role, which should also frame them as heroes—a nuance that was completely disregarded. The DeepSeek models provide no reasoning for this misclassification.

One possible explanation is that the sentence structure was too complex for the models to properly connect "their role" back to "women". This suggests that either a model bias or unclear prompt instructions could be responsible, but it remains uncertain which played a greater role.

The following example illustrates why relying solely on LLMs for tasks like character framing is dangerous:

- (19) "[...] the adoption of a plan of action to combat sexual violence among **the armed forces**

[...] (UNSC_2016_SPV.7704_spch087.txt)

In this case, "the armed forces" should clearly be labeled as villains, as they are described as perpetrators of sexual violence. LlamaZero and LlamaFew assigned the correct label but made a structural mistake by annotating only "armed forces", omitting "the", and thus failing to comply with annotation guidelines.

However, the errors made by DeepSeekZero and DeepSeekFew are far more concerning. DeepSeekZero failed to annotate the entity at all, while DeepSeekFew went so far as to label "the armed forces" as heroes. The reasoning provided by DeepSeekFew sheds some light on why this might have happened:

"The Government of the DRC is taking actions like appointing representatives, adopting plans, and signing decrees. These are proactive steps to combat sexual violence, so the Government is a Hero. Specific roles like the personal representative of the President, the Minister of National Defence, and the military justice system are Heroes as they're part of the government's efforts."

Since the plan to combat sexual violence among the armed forces was introduced by the government, the model may have mistakenly associated "the armed forces" with the military justice system—which is part of the government's anti-violence efforts—thus framing them positively. DeepSeekZero, on the other hand, offers no explanation at all for its omission.

This suggests that the model struggles to understand the full context of the text and may overlook even explicit accusations—though in this case, the accusation is anything but subtle.

In summary, the error analysis reveals that while some models, like RoBERTa, generally perform well in terms of entity recognition, challenges remain in handling complex cases involving articles, pronouns, multi-word entities, and contextual role shifts. The few-shot prompting approaches showed some improvement, particularly in models like DeepSeekFew and LlamaFew, but inconsistencies and biases still persist. Further refinements and understanding of these models' behavior, especially in handling gender-related and multi-label scenarios, are crucial for improving their performance in annotation tasks.

5 Conclusion

Inspired by the central question of how women are depicted in a dataset that revolves around them, this paper conducted a thorough analysis, with a particular focus on their roles as heroes, victims, and villains. Using topic modeling, clustering, and manual annotations of 54 speeches, we confirmed findings from previous work related to the WPS agenda, especially concerning women. For instance, we found that women are depicted as the main characters and the primary victims in open debates. However, they also appear as agents of change, providing valuable knowledge and playing a significant role in peacekeeping missions. As victims, they often co-occur with girls and children, reinforcing their portrayal as symbols of innocent civilians. They are never depicted as villains—only within descriptions of villains committing harm against them.

Men, on the other hand, are mentioned far less frequently. When they are, it is either alone as villains or, more commonly, in combination with women, children, and boys—either as victims or as heroes alongside women, emphasizing the importance of cooperation between genders. Meanwhile, the Security Council and its members primarily portray themselves as heroes by highlighting their contributions, though occasionally they acknowledge their own shortcomings, subtly villainizing themselves by admitting they are not doing enough.

Other than that, villains are rarely mentioned, and when they are, they are described in vague terms. While some non-state actors, such as ISIS, are explicitly called out, most of the time, perpetrators remain unnamed, and no member state is directly accused. This reflects the diplomatic nature of UN discourse, where positive, optimistic visions are encouraged.

Furthermore, we explored which framework could enable the automatic identification of roles for occurring entities. To do so, we annotated 54 speeches from the WPS corpus. We then tested the LLM RoBERTa as well as chat-optimized LLMs such as DeepSeek and Llama, using zero- and few-shot prompting. The results showed that the chat-optimized LLMs performed surprisingly poorly, often ignoring key rules from the annotation guidelines and even misclassifying obvious labels—for instance, labeling villains as heroes. RoBERTa, on the other hand, proved to be quite reliable and was even able to identify some entities that the annota-

tors had missed.

However, especially in the context of gender and institutional entities, we cannot exclude that the models introduce their own biases in the annotation process. Interestingly, DeepSeek’s reasoning feature at least provided explanations for some of its misclassifications, offering valuable insights into the black box of LLM decision-making. While there is still significant room for improvement, these findings offer a strong first step toward automated entity role identification.

Limitations

Annotations One of the primary limitations of this study is the scope of our annotations. Given the constraints of the project time frame and the fact that we were a small team balancing other responsibilities, it was not feasible to annotate a larger number of speeches. A more extensive dataset would have provided additional training data and enabled a more robust evaluation of our approaches. While the annotation process itself was not overly time-consuming, significant effort was required to define the annotation guidelines and curate the speeches collaboratively, making this stage particularly resource-intensive.

Moreover, character roles are inherently subjective, meaning that we, as annotators, inevitably brought our own biases into the annotation process. This likely influenced both the evaluation and the model’s performance, as the training data was shaped by our individual interpretations. These subjective factors further highlight the challenges of automating entity role classification and the need for continuous refinement of guidelines and methodologies.

Frameworks Another limitation pertains to the computational and accessibility constraints associated with the models used. Many cutting-edge LLMs are not freely available and require subscription-based access, limiting their usability for broader research applications. Additionally, fine-tuning models such as RoBERTa demands substantial computational resources, which were restricted in our case since we relied on personal devices. As a result, we were limited in the extent of hyperparameter tuning we could perform. Interestingly, the relatively small size of our annotated dataset was a double-edged sword, as it reduced the computational burden, making some experiments feasible, but also limited the overall robustness of

our findings.

Future Work

Our analysis has provided valuable insights into the depiction of women and men in the WPS speeches. However, several questions remain open for further investigation, such as examining differences in framing across countries or time periods. Additionally, our frameworks for automatic detection present opportunities for refinement and comparative analysis. Below, we outline several promising directions for future research.

Addressing Class Imbalance Future work should aim to mitigate class imbalance by augmenting the dataset with additional examples of underrepresented roles, particularly in multi-label classification scenarios. This would help ensure that all role categories are sufficiently represented, improving model performance and robustness.

Comparing with and/or Integrating Previous Methods While our approach was developed as a stand-alone framework, it would be valuable to apply existing methodologies, such as those proposed by [Stammach et al. \(2022\)](#) and [Gomez-Zara et al. \(2018a\)](#), to our dataset. A comparative analysis could help identify strengths and weaknesses in different annotation strategies. Furthermore, integrating elements from these prior approaches into our framework could enhance the automatic detection process.

References

- Çağlayan Başer. 2024. Women’s role in violence and un women, peace, and security agenda. *Alternatif Politika*, 16(1):1–30.
- Kelly Bergstrand and James M. Jasper. 2018. [Villains, victims, and heroes in character theory and affect control theory](#). *Social Psychology Quarterly*, 81(3):228–247.
- R. Charli Carpenter. 2005. [“women, children and other vulnerable groups”](#): Gender, strategic frames and the protection of civilians as a transnational issue. *International Studies Quarterly*, 49(2):295–334.
- Stanley Cohen. 2011. *Folk devils and moral panics*. Routledge.
- Victoria Martín de la Rosa and Luis Miguel Lázaro. 2019. [How women are imagined through conceptual metaphors in united nations security council resolutions on women, peace and security](#). *Journal of Gender Studies*, 28(4):373–386.

- Martin Ester, Hans-Peter Kriegel, Jörg Sander, Xiaowei Xu, et al. 1996. A density-based algorithm for discovering clusters in large spatial databases with noise. In *kdd*, volume 96, pages 226–231.
- Lea Frermann, Jiatong Li, Shima Khanehazar, and Gosia Mikolajczak. 2023. Conflicts, villains, resolutions: Towards models of narrative media framing. *arXiv preprint arXiv:2306.02052*.
- Kai Gehring and Matteo Grigoletto. 2023. Analyzing climate change policy narratives with the character-role narrative framework. Technical report, CESifo Working Paper.
- Aurélien Geron. 2022. *Hands-on machine learning with Scikit-Learn, Keras, and TensorFlow: Concepts, tools, and techniques to build intelligent systems*. "O'Reilly Media, Inc."
- Sheri Lynn Gibbings. 2011. [No angry women at the united nations: Political dreams and the cultural politics of united nations security council resolution 1325](#). *International Feminist Journal of Politics*, 13(4):522–538.
- Diego Gomez-Zara, Miriam Boon, and Larry Birnbaum. 2018a. Who is the hero, the villain, and the victim? detection of roles in news articles using natural language techniques. In *Proceedings of the 23rd International Conference on Intelligent User Interfaces*, pages 311–315.
- Diego Gomez-Zara, Miriam Boon, and Larry Birnbaum. 2018b. [Who is the hero, the villain, and the victim? detection of roles in news articles using natural language techniques](#). In *Proceedings of the 23rd International Conference on Intelligent User Interfaces*, IUI '18, page 311–315, New York, NY, USA. Association for Computing Machinery.
- Aaron Grattafiori, Abhimanyu Dubey, Abhinav Jauhri, Abhinav Pandey, Abhishek Kadian, Ahmad Al-Dahle, Aiesha Letman, Akhil Mathur, Alan Schelten, Alex Vaughan, Amy Yang, Angela Fan, Anirudh Goyal, Anthony Hartshorn, Aobo Yang, Archi Mitra, Archie Sravankumar, Artem Korenev, Arthur Hinsvark, Arun Rao, Aston Zhang, Aurelien Rodriguez, Austen Gregerson, Ava Spataru, Baptiste Roziere, Bethany Biron, Binh Tang, Bobbie Chern, Charlotte Caucheteux, Chaya Nayak, Chloe Bi, Chris Marra, Chris McConnell, Christian Keller, Christophe Touret, Chunyang Wu, Corinne Wong, Cristian Canton Ferrer, Cyrus Nikolaidis, Damien Al-lonsius, Daniel Song, Danielle Pintz, Danny Livshits, Danny Wyatt, David Esiobu, Dhruv Choudhary, Dhruv Mahajan, Diego Garcia-Olano, Diego Perino, Dieuwke Hupkes, Egor Lakomkin, Ehab AlBadawy, Elina Lobanova, Emily Dinan, Eric Michael Smith, Filip Radenovic, Francisco Guzmán, Frank Zhang, Gabriel Synnaeve, Gabrielle Lee, Georgia Lewis Anderson, Govind Thattai, Graeme Nail, Gregoire Milon, Guan Pang, Guillem Cucurell, Hailey Nguyen, Hannah Korevaar, Hu Xu, Hugo Touvron, Iliyan Zarov, Imanol Arrieta Ibarra, Isabel Kloumann, Ishan Misra, Ivan Evtimov, Jack Zhang, Jade Copet, Jaewon Lee, Jan Geffert, Jana Vranes, Jason Park, Jay Mahadeokar, Jeet Shah, Jelmer van der Linde, Jennifer Billock, Jenny Hong, Jenya Lee, Jeremy Fu, Jianfeng Chi, Jianyu Huang, Jiawen Liu, Jie Wang, Jiecao Yu, Joanna Bitton, Joe Spisak, Jongsoo Park, Joseph Rocca, Joshua Johnstun, Joshua Saxe, Junteng Jia, Kalyan Vasuden Alwala, Karthik Prasad, Kartikeya Upasani, Kate Plawiak, Ke Li, Kenneth Heafield, Kevin Stone, Khalid El-Arini, Krithika Iyer, Kshitiz Malik, Kuenley Chiu, Kunal Bhalla, Kushal Lakhotia, Lauren Rantala-Yeary, Laurens van der Maaten, Lawrence Chen, Liang Tan, Liz Jenkins, Louis Martin, Lovish Madaan, Lubo Malo, Lukas Blecher, Lukas Landzaat, Luke de Oliveira, Madeline Muzzi, Mahesh Pasupuleti, Mannat Singh, Manohar Paluri, Marcin Kardas, Maria Tsimpoukelli, Mathew Oldham, Mathieu Rita, Maya Pavlova, Melanie Kam-badur, Mike Lewis, Min Si, Mitesh Kumar Singh, Mona Hassan, Naman Goyal, Narjes Torabi, Nikolay Bashlykov, Nikolay Bogoychev, Niladri Chatterji, Ning Zhang, Olivier Duchenne, Onur Çelebi, Patrick Alrassy, Pengchuan Zhang, Pengwei Li, Petar Vasic, Peter Weng, Prajjwal Bhargava, Pratik Dubal, Praveen Krishnan, Punit Singh Koura, Puxin Xu, Qing He, Qingxiao Dong, Ragavan Srinivasan, Raj Ganapathy, Ramon Calderer, Ricardo Silveira Cabral, Robert Stojnic, Roberta Raileanu, Rohan Maheswari, Rohit Girdhar, Rohit Patel, Romain Sauvestre, Ronnie Polidoro, Roshan Sumbaly, Ross Taylor, Ruan Silva, Rui Hou, Rui Wang, Saghar Hosseini, Sahana Chennabasappa, Sanjay Singh, Sean Bell, Seohyun Sonia Kim, Sergey Edunov, Shao-liang Nie, Sharan Narang, Sharath Raparthy, Sheng Shen, Shengye Wan, Shruti Bhosale, Shun Zhang, Simon Vandenhende, Soumya Batra, Spencer Whitman, Sten Sootla, Stephane Collot, Suchin Gururangan, Sydney Borodinsky, Tamar Herman, Tara Fowler, Tarek Sheasha, Thomas Georgiou, Thomas Scialom, Tobias Speckbacher, Todor Mihaylov, Tong Xiao, Ujjwal Karn, Vedanuj Goswami, Vibhor Gupta, Vignesh Ramanathan, Viktor Kerkez, Vincent Gonguet, Virginie Do, Vish Vogeti, Vitor Albiero, Vladan Petrovic, Weiwei Chu, Wenhan Xiong, Wenyin Fu, Whitney Meers, Xavier Martinet, Xiaodong Wang, Xiaofang Wang, Xiaoqing Ellen Tan, Xide Xia, Xinfeng Xie, Xuchao Jia, Xuwei Wang, Yaelle Goldschlag, Yashesh Gaur, Yasmine Babaei, Yi Wen, Yiwen Song, Yuchen Zhang, Yue Li, Yuning Mao, Zacharie Del-pierre Coudert, Zheng Yan, Zhengxing Chen, Zoe Papakipos, Aaditya Singh, Aayushi Srivastava, Abha Jain, Adam Kelsey, Adam Shajnfeld, Adithya Gangidi, Adolfo Victoria, Ahuva Goldstand, Ajay Menon, Ajay Sharma, Alex Boesenberg, Alexei Baevski, Allie Feinstein, Amanda Kallet, Amit Sangani, Amos Teo, Anam Yunus, Andrei Lupu, Andres Alvarado, Andrew Caples, Andrew Gu, Andrew Ho, Andrew Poulton, Andrew Ryan, Ankit Ramchandani, Annie Dong, Annie Franco, Anuj Goyal, Aparajita Saraf, Arkabandhu Chowdhury, Ashley Gabriel, Ashwin Bharambe, Assaf Eisenman, Azadeh Yazdan, Beau James, Ben Maurer, Benjamin Leonhardi, Bernie Huang, Beth Loyd, Beto De Paola, Bhargavi Paranjape, Bing Liu, Bo Wu, Boyu Ni, Braden Hancock, Bram Wasti, Brandon Spence, Brani Stojkovic,

- Brian Gamido, Britt Montalvo, Carl Parker, Carly Burton, Catalina Mejia, Ce Liu, Changan Wang, Changkyu Kim, Chao Zhou, Chester Hu, Ching-Hsiang Chu, Chris Cai, Chris Tindal, Christoph Feichtenhofer, Cynthia Gao, Damon Civin, Dana Beaty, Daniel Kreymer, Daniel Li, David Adkins, David Xu, Davide Testuggine, Delia David, Devi Parikh, Diana Liskovich, Didem Foss, Dingkan Wang, Duc Le, Dustin Holland, Edward Dowling, Eissa Jamil, Elaine Montgomery, Eleonora Presani, Emily Hahn, Emily Wood, Eric-Tuan Le, Erik Brinkman, Esteban Arcaute, Evan Dunbar, Evan Smothers, Fei Sun, Felix Kreuk, Feng Tian, Filippas Kokkinos, Firat Ozgenel, Francesco Caggioni, Frank Kanayet, Frank Seide, Gabriela Medina Florez, Gabriella Schwarz, Gada Badeer, Georgia Swee, Gil Halpern, Grant Herman, Grigory Sizov, Guangyi, Zhang, Guna Lakshminarayanan, Hakan Inan, Hamid Shojanazeri, Han Zou, Hannah Wang, Hanwen Zha, Haroun Habeeb, Harrison Rudolph, Helen Suk, Henry Aspegren, Hunter Goldman, Hongyuan Zhan, Ibrahim Damlaj, Igor Molybog, Igor Tufanov, Ilias Leontiadis, Irina-Elena Veliche, Itai Gat, Jake Weissman, James Geboski, James Kohli, Janice Lam, Japhet Asher, Jean-Baptiste Gaya, Jeff Marcus, Jeff Tang, Jennifer Chan, Jenny Zhen, Jeremy Reizenstein, Jeremy Teboul, Jessica Zhong, Jian Jin, Jingyi Yang, Joe Cummings, Jon Carvill, Jon Shepard, Jonathan McPhie, Jonathan Torres, Josh Ginsburg, Junjie Wang, Kai Wu, Kam Hou U, Karan Saxena, Kartikay Khandelwal, Katayoun Zand, Kathy Matosich, Kaushik Veeraraghavan, Kelly Michelen, Keqian Li, Kiran Jagadeesh, Kun Huang, Kunal Chawla, Kyle Huang, Lailin Chen, Lakshya Garg, Lavender A, Leandro Silva, Lee Bell, Lei Zhang, Liangpeng Guo, Licheng Yu, Liron Moshkovich, Luca Wehrstedt, Madian Khabza, Manav Avalani, Manish Bhatt, Martynas Mankus, Matan Hasson, Matthew Lennie, Matthias Reso, Maxim Groshev, Maxim Naumov, Maya Lathi, Meghan Keneally, Miao Liu, Michael L. Seltzer, Michal Valko, Michelle Restrepo, Mihir Patel, Mik Vyatskov, Mikayel Samvelyan, Mike Clark, Mike Macey, Mike Wang, Miquel Jubert Hermoso, Mo Metanat, Mohammad Rastegari, Munish Bansal, Nandhini Santhanam, Natascha Parks, Natasha White, Navyata Bawa, Nayan Singhal, Nick Egebo, Nicolas Usunier, Nikhil Mehta, Nikolay Pavlovich Laptev, Ning Dong, Norman Cheng, Oleg Chernoguz, Olivia Hart, Omkar Salpekar, Ozlem Kalinli, Parkin Kent, Parth Parekh, Paul Saab, Pavan Balaji, Pedro Rittner, Philip Bontrager, Pierre Roux, Piotr Dollar, Polina Zvyagina, Prashant Ratanchandani, Pritish Yuvraj, Qian Liang, Rachad Alao, Rachel Rodriguez, Rafi Ayub, Raghotham Murthy, Raghu Nayani, Rahul Mitra, Rangaprabhu Parthasarathy, Raymond Li, Rebekkah Hogan, Robin Battey, Rocky Wang, Russ Howes, Ruty Rinott, Sachin Mehta, Sachin Siby, Sai Jayesh Bondu, Samyak Datta, Sara Chugh, Sara Hunt, Sargun Dhillon, Sasha Sidorov, Satadru Pan, Saurabh Mahajan, Saurabh Verma, Seiji Yamamoto, Sharadh Ramaswamy, Shaun Lindsay, Shaun Lindsay, Sheng Feng, Shenghao Lin, Shengxin Cindy Zha, Shishir Patil, Shiva Shankar, Shuqiang Zhang, Shuqiang Zhang, Sinong Wang, Sneha Agarwal, Soji Sajuyigbe, Soumith Chintala, Stephanie Max, Stephen Chen, Steve Kehoe, Steve Satterfield, Sudarshan Govindaprasad, Sumit Gupta, Summer Deng, Sungmin Cho, Sunny Virk, Suraj Subramanian, Sy Choudhury, Sydney Goldman, Tal Remez, Tamar Glaser, Tamara Best, Thilo Koehler, Thomas Robinson, Tianhe Li, Tianjun Zhang, Tim Matthews, Timothy Chou, Tzook Shaked, Varun Vontimitta, Victoria Ajayi, Victoria Montanez, Vijai Mohan, Vinay Satish Kumar, Vishal Mangla, Vlad Ionescu, Vlad Poenaru, Vlad Tiberiu Mihailescu, Vladimir Ivanov, Wei Li, Wenchen Wang, Wenwen Jiang, Wes Bouaziz, Will Constable, Xiaocheng Tang, Xiaojuan Wu, Xiaolan Wang, Xilun Wu, Xinbo Gao, Yaniv Kleinman, Yanjun Chen, Ye Hu, Ye Jia, Ye Qi, Yenda Li, Yilin Zhang, Ying Zhang, Yossi Adi, Youngjin Nam, Yu, Wang, Yu Zhao, Yuchen Hao, Yundi Qian, Yunlu Li, Yuzi He, Zach Rait, Zachary DeVito, Zef Rosnbrick, Zhaoduo Wen, Zhenyu Yang, Zhiwei Zhao, and Zhiyu Ma. 2024. [The llama 3 herd of models](#). *Preprint*, arXiv:2407.21783.
- Maarten Grootendorst. 2022. Bertopic: Neural topic modeling with a class-based tf-idf procedure. *arXiv preprint arXiv:2203.05794*.
- Daya Guo, Dejian Yang, Haowei Zhang, Junxiao Song, Ruoyu Zhang, Runxin Xu, Qihao Zhu, Shirong Ma, Peiyi Wang, Xiao Bi, et al. 2025. Deepseek-r1: Incentivizing reasoning capability in llms via reinforcement learning. *arXiv preprint arXiv:2501.12948*.
- James M Jasper, Michael Young, and Elke Zuern. 2018. Character work in social movements. *Theory and Society*, 47:113–131.
- Stephen Karpman. 1968. Fairy tales and script drama analysis. *Transactional analysis bulletin*, 7(26):39–43.
- Orrin E Klapp. 1954. Heroes, villains and fools, as agents of social control. *American sociological review*, 19(1):56–62.
- Jan-Christoph Klie, Michael Bugert, Beto Boullosa, Richard Eckart De Castilho, and Iryna Gurevych. 2018. The inception platform: Machine-assisted and knowledge-oriented interactive annotation. In *Proceedings of the 27th international conference on computational linguistics: system demonstrations*, pages 5–9.
- Ludovic Kun, Jayesh Bankoti, and David Kiskovsk. 2022. [Logically at the constraint 2022: Multimodal role labelling](#). In *Proceedings of the Workshop on Combating Online Hostile Posts in Regional Languages during Emergency Situations*, pages 24–34, Dublin, Ireland. Association for Computational Linguistics.
- Xianming Li and Jing Li. 2023. Angle-optimized text embeddings. *arXiv preprint arXiv:2309.12871*.
- Yinhan Liu, Myle Ott, Naman Goyal, Jingfei Du, Mandar Joshi, Danqi Chen, Omer Levy, Mike Lewis, Luke Zettlemoyer, and Veselin Stoyanov. 2019.

- Roberta: A robustly optimized bert pretraining approach. *Preprint*, arXiv:1907.11692.
- Leland McInnes, John Healy, and James Melville. 2018. Umap: Uniform manifold approximation and projection for dimension reduction. *arXiv preprint arXiv:1802.03426*.
- Hiroki Nakayama. 2018. [sequeval: A python framework for sequence labeling evaluation](https://github.com/chakki-works/sequeval). Software available from <https://github.com/chakki-works/sequeval>.
- Hannah Rashkin, Sameer Singh, and Yejin Choi. 2015. Connotation frames: A data-driven investigation. *arXiv preprint arXiv:1506.02739*.
- Maarten Sap, Marcella Cindy Prasettio, Ari Holtzman, Hannah Rashkin, and Yejin Choi. 2017. Connotation frames of power and agency in modern films. In *Proceedings of the 2017 conference on empirical methods in natural language processing*, pages 2329–2334.
- Mirco Schoenfeld, Steffen Eckhard, Ronny Patz, Hilde van Meegdenburg, and Antonio Pires. 2019. [The UN Security Council Debates](#).
- Erich Schubert, Jörg Sander, Martin Ester, Hans Peter Kriegel, and Xiaowei Xu. 2017. DbSCAN revisited, revisited: why and how you should (still) use dbSCAN. *ACM Transactions on Database Systems (TODS)*, 42(3):1–21.
- Shivam Sharma, Atharva Kulkarni, Tharun Suresh, Himanshi Mathur, Preslav Nakov, Md. Shad Akhtar, and Tanmoy Chakraborty. 2023. [Characterizing the entities in harmful memes: Who is the hero, the villain, the victim?](#) *Preprint*, arXiv:2301.11219.
- Shivam Sharma, Tharun Suresh, Atharva Kulkarni, Himanshi Mathur, Preslav Nakov, Md. Shad Akhtar, and Tanmoy Chakraborty. 2022. [Findings of the CONSTRAINT 2022 shared task on detecting the hero, the villain, and the victim in memes](#). In *Proceedings of the Workshop on Combating Online Hostile Posts in Regional Languages during Emergency Situations*, pages 1–11, Dublin, Ireland. Association for Computational Linguistics.
- Pranaydeep Singh, Aaron Maladry, and Els Lefever. 2022. Combining language models and linguistic information to label entities in memes. In *Proceedings of the Workshop on Combating Online Hostile Posts in Regional Languages during Emergency Situations*, pages 35–42.
- Dominik Stammach, Maria Antoniak, and Elliott Ash. 2022. [Heroes, villains, and victims, and GPT-3: Automated extraction of character roles without training data](#). In *Proceedings of the 4th Workshop of Narrative Understanding (WNU2022)*, pages 47–56, Seattle, United States. Association for Computational Linguistics.
- Erik F. Tjong Kim Sang and Sabine Buchholz. 2000. [Introduction to the CoNLL-2000 shared task chunking](#). In *Fourth Conference on Computational Natural Language Learning and the Second Learning Language in Logic Workshop*.
- Jacqui True and Antje Wiener. 2019. [Everyone wants \(a\) peace: the dynamics of rhetoric and practice on ‘women, peace and security’](#). *International Affairs*, 95(3):553–574.
- UNDP. 2023. [Parliament as partners: Supporting the women, peace, and security agenda - a global handbook](#). Accessed: 2025-03-24.
- Yue Wu, Xuan Tang, Tom M. Mitchell, and Yuanzhi Li. 2024. [Smartplay: A benchmark for llms as intelligent agents](#). *Preprint*, arXiv:2310.01557.

Appendix A Plots for Annotation Data Results




Figure 18: Distribution of Characters for Women Entities

Mentions of Entities Containing ['men', "men's"] Across Categories




Figure 19: Distribution of Characters for Men Entities

Appendix B Interannotator Agreement

	Annotator 1	Annotator 2	Annotator 3	Curator
Annotator 1	–	0.30	0.44	0.47
Annotator 2		–	0.38	0.48
Annotator 3			–	0.70
Curator				–

Table 7: Krippendorff’s Alpha (unitizing/character offsets) for agreement on role annotation (Hero, Victim, Villain) computed pairwise. We three annotators worked collaboratively to define a gold standard, which is reflected in the Curator column.

Appendix C Heroes, Victims, and Perpetrators in WPS: Annotation Guidelines

i. Document Level Rules Exclude all president speeches

ii. Character Level Rules

a. We include entire noun phrases (NP): This includes any numerical or descriptive modifiers.

Example: “60 million Africans”

b. We annotate restrictive relative clauses fully.

Example: “parties that perpetuate acts of violence against women and children”

c. We include possessive modifiers within the NP:

Example: “my delegation”

d. When multiple characters are listed together as part of the same NP and share a common role or action, we consider the entire sequence as one entity, so we not split or annotate each item separately.

Example: “Allow me to begin by thanking **the Secretary-General, Mr. Kevin Hyland, Mr. Yury Fedotov and Ms. Ilwad Elman** for their briefings.”

e. We do not include punctuation at the end of annotations.

Example: “**men and boys.**”

f. We do not annotate predicate nominatives or any other descriptions of characters.

Example: “**Women** are not only victims.” Here, ‘victims’ should not be annotated.

Example: “**women** as agents of peace”

g. We do not annotate positions or groups when they are only mentioned in the abstract, without references to the achievements of the particular entity.

Example: “the establishment of the post of the Special Representative of the Secretary General on sexual violence in situations of armed conflict”

Role-Definitions

In German, the constellation of hero, victim, and perpetrator is called the drama triangle, first introduced by [Karpman \(1968\)](#). It is a psychological model used in transactional analysis to describe interactions between individuals or groups as adopting one of three stereotypical roles: victim, perpetrator, or hero, which simplifies the linking of behaviors with stereotypical roles learned at an early stage ([Karpman, 1968](#); [Gomez-Zara et al., 2018a](#)). The link to stereotypes, the framing of the characters with said roles, influences the readers to have a certain bias and to develop a non-reflective opinion through the stereotypical knowledge, which is influenced by the author alone ([Gomez-Zara et al., 2018a](#)). Hero, victim, and perpetrator are also of great importance in affect control theory and character theory ([Bergstrand and Jasper, 2018](#)). The following are definitions we will work with for the WPS-project:

Victim

Victims tend to be portrayed as weak, good, innocent people who are in need of protection ([Jasper et al., 2018](#)). Due to these characteristics, they often motivate and encourage action towards a specific cause and can help make aware of injustices which are worth combating. [Jasper et al. \(2018\)](#) state that victim’s sufferings are often elaborated in detail to arouse more moral emotions and indignation. “Popular” victims, as they get the most sympathetic reactions in the modern world due to their cultural innocence, are children ([Jasper et al., 2018](#)).

Task-specific additions (based on the WPS corpus):

- Someone who is excluded from decision-making processes/someone who is not given the recognition/power that they deserve.
- Someone who suffers acts of sexual violence/physical harm/displacement, etc.
- Someone who is not given the same equal rights as other parties

Perpetrator/Villain

Perpetrators are people whose moral reputation turns or has turned negative (Jasper et al., 2018). They are considered to be people who spread anxiety and fear, cause people to lose their daily routines, and make them sacrifice their lives, for example, within wars (Bergstrand and Jasper, 2018). Perpetrators often share the same characteristics as heroes, such as being strong, brave, and intelligent. However, their description tends to be more like that of beast-like predators: powerful, threatening, and delinquent (Cohen, 2011).

Task specific additions (based on the WPS corpus):

- Someone who is responsible for causing anxieties, damage, and crimes.
- Someone who stands in the way of equal rights and justice for victims.

Hero

Heroes are people who, by helping victims (and defeating the villains (Jasper et al., 2018; Klapp, 1954)), can become heroes. They are defined as people who increase agreement within groups and boost commitment to a cause. They tend to be well-intentioned people, who recognize injustice, try to resolve and fight it, as well as protect others. However, this does not mean that heroes are completely independent. Jasper et al. (2018) state that even a hero might be in need of help from an even more experienced hero. Furthermore, they are often put in the context of success (Klapp, 1954).

Task specific additions (based on the WPS corpus):

- Someone who has the potential to save others, given their abilities, knowledge, positionality/perspective, but who is not able to do so because of structural discrimination.
- Someone who makes a call to action
- Someone who is recognizing the unjust treatment and violation of victims or is calling it out.

Further Thoughts

It can be seen that the victim, hero, and perpetrator definitions align very much with stereotypes we are taught and confronted with since our childhood. However, roles are not finite, as heroes can turn into victims or perpetrators, and just like that, all the other roles can take on another one (Jasper et al., 2018). The definitions are tried to be kept as simple and generally adaptive as possible, because depending on the source one works with, representations can change. Jasper et al. (2018), for example, state that in news articles, victims are being kept nameless in contrast to heroes and perpetrators. In the WPS dataset, however, we observed that it is the perpetrators who are mainly being kept nameless and victims who are being mentioned frequently—still nameless beings, but specified as “women”, “children”, etc.

Entities to tag

- a. People
 - b. Organisations
 - c. Countries
 - d. Groups
 - e. UN Resolutions. Only annotate United Nations Resolutions if they are explicitly personified.
- Example:* In “Resolution 1325 calls for action,” the resolution should be tagged as a hero. In “working towards the implementation of resolution 1325,” the resolution should not be marked.

Entities NOT to tag

- a. Abstract concepts: Abstract ideas or symbolic references.
Examples: “International cooperation,” “sexual violence,” “**women’s** participation.” In the latter example, women are the entities that are taking heroic action, while their participation refers to a concept.
- b. Entities hoping for/welcoming/thanking/commending something good. These are passive actions.

c. Laws or Treaties.

Annotation of character role labels

We only annotate generic role terms (like “victim”) when no other specification is included.

Example: In “Victims of these atrocious crimes have been waiting for justice”, we would tag “victims of these atrocious crimes” as Victim. However, if specific entities are mentioned, like in “Victims of these atrocious crimes, namely women, have been waiting for justice” we only annotate “women” as Victim.

Annotation of multiple roles

In certain cases, characters may be portrayed with multiple roles simultaneously. When this happens, we annotate the entity with a combined role. Furthermore, the same entity can take on different roles throughout the speech. For instance, while “women” might be classified as ‘victim’ in one sentence, the same entity can also be classified as ‘hero’ in a different part of the speech.

Example: “The equal right to decision-making and participation, along with **women’s** empowerment, is crucial to ensure a functioning society and peace and justice in the aftermath of conflicts.”

Explanation: In this case, “women” are portrayed both as victims (since they need external help to be empowered) as well as heroes (since they contribute to a peaceful society). When an entity fits into multiple roles based on the context, we use combined tags.

iii. Annotation examples and explanations

Legend: Hero Perpetrator Victim

1. "We commend the work that has been done by the United Nations Children’s Fund in reintegration projects that has led to the release of girls from the armed forces in various countries."

In this example, "we" is not annotated as a hero because they are just thanking another organization. However, the organization in question, the "United Nations Children’s Fund", is marked as a hero because they are responsible for these projects that help girls. The "girls", on the other hand, are marked as victims because, although it could be argued that they are the benefactors of the situation (since they have been released), we believe that the fact that they were abducted in the first place carries more weight, also because they are in danger of being abducted again at any time. Most intuitively, "the armed forces" are the perpetrators, since they were responsible for the girls’ abduction.

2. "Members of the Council note that although women have begun to play an important role in conflict resolution, peacekeeping and peace- building, they are still under-represented in decision-making in regard to conflict. If women are to play an equal part in security and maintaining peace, they must be empowered politically and economically, and represented adequately at all levels of decision-making, both at the pre-conflict stage and during hostilities, as well as at the point of peacekeeping, peace-building, reconciliation and reconstruction."

This example demonstrates that the same entity (in this case, “women”) can take on multiple roles within one sentence. In the first instance (“women”), they are marked as heroes due to their important role in conflict resolution. However, in the second part of the sentence, where they are referred to as “they,” the focus shifts to their underrepresentation in decision-making, marking them as victims. The conjunction "although" introduces this contrast, highlighting both aspects of the situation. In the second sentence, the focus remains on women as victims (“if women are to play an equal part in security”), implying that they are not currently able to do so. The need for empowerment, both politically and economically, is presented because they are not given equal opportunities as men. Therefore, they are portrayed as victims in this context.

3. "Those who commit crimes against women, including the peacekeeping personnel, should be brought to book. Let us heed the women's cry for an equal opportunity to voice their ideas in official peace negotiations. And let us act now."

In the last two sentences of this example, "us" is marked as a hero since the speaker is engaging in a call to action. The speaker portrays "us" as a hero by attempting to encourage others to acknowledge and act upon the unjust inequality faced by women.

4. "My Government attaches great importance to the protection and security of women and girls, both in situations of armed conflict and in peace."

In this case, "my government" is marked as a hero because, while the action is indirect, they recognize the importance of the issue and stand up for it. "Women and girls" are marked as victims since they require protection and security, as they cannot protect themselves (see victim definition above). Although they benefit from the protection measures (and thus could be seen as benefactors), we chose to prioritize their need for protection, which highlights their vulnerable position and role as victims in society.

5. "Being direct victims of violence and discrimination, women have gained a great understanding of the need to address peace comprehensively."

Here, "women" are simultaneously considered heroes and victims. The victim annotation is indicated by the context of them being the "direct victim," while the hero annotation relies on the positionality of women—their potential to provide valuable knowledge.

Appendix D Hyperparameters

	Learning Rate	Epochs	Optimizer
m1	5e-5	3	AdamW
m2	5e-5	30	AdamW
m3	1e-5	3	AdamW
m4	1e-5	30	AdamW
m5	5e-5	30	SGD
m6	3e-5	30	SGD

Appendix E Zero-shot prompt

Your goal: Rewrite the input speech with only in-line character role annotations, using precise tags, leaving spacing and spelling as is, and following all structural and role-based rules.

Task overview You are given a speech from the United Nations Security Council. Your task is to identify and label characters within the speech. Label each identified character as either Hero, Villain, or Victim. These labels can also be combined to mark that an identified character is portrayed as fulfilling several roles, for example, as Hero and Victim. Some character spans may contain other character spans, for example, when a victim entity is mentioned within a villain entity.

For the annotations, follow these annotation guidelines:

Annotation guidelines

1. Character Level Rules:

a. Include entire noun phrases (NP): This includes any numerical or descriptive modifiers. Example: “<HER>60 million Africans</HER> ”

b. Annotate restrictive relative clauses fully. Example: “<HER>parties that perpetuate acts of violence against women and children</HER> ”

c. Include possessive modifiers within the NP:

Example: “<HER>my delegation</HER> ”

d. When multiple characters are listed together as part of the same NP and share a common role or action, tag the entire sequence as one entity. Do not split or annotate each item separately. Example: “Allow me to begin by thanking <HER>the Secretary-General, Mr. Kevin Hyland, Mr. Yury Fedotov and Ms. Ilwad Elman</HER> for their briefings.”

e. Do not include punctuation at the end of annotations. Example: “<VIC>men and boys</VIC>.” f. Do not annotate predicate nominatives or any other descriptions of characters.

Example: “<VIC>Women</VIC> are not only victims.” Here, victims should not be annotated.

Example: “<HER>women</HER> as agents of peace”

g. Do not annotate positions or groups when they are only mentioned in the abstract, without references to the achievements of the particular entity.

Example: “the establishment of the post of the Special Representative of the Secretary General on sexual violence in situations of armed conflict”

2. Role-Definitions:

a. Hero: An entity is labeled as a Hero if it: helps victims and/or defeats villains increases agreement within groups or commitment to a cause recognizes injustice and actively fights to resolve it has the potential to save others, given their abilities, knowledge, positionality/perspective, but who is not able to do so because of structural discrimination makes a call to action calls out or recognizes unjust treatment/violations

b. Victim: An entity is labeled as a Victim if it: is weak, good, or innocent and in need of protection suffers from injustice (e. g., sexual violence, displacement, physical harm) is excluded from decision-making or denied recognition/power does not have the same rights as others

c. Villain: An entity is labeled as a Villain if it: causes anxiety and fear causes people to lose their daily routines causes harm or disrupts peace acts in a way that prevents equal rights and justice for victims has a negative moral reputation

3. Entities to tag: You should only tag:

a. People

b. Organisations

c. Countries

d. Groups

e. UN Resolutions. Only annotate United Nations Resolutions if they are explicitly personified.

Example: In “Resolution 1325 calls for action,” the resolution should be tagged as a hero. In “working towards the implementation of resolution 1325,” the resolution should not be marked.

4. Entities NOT to tag:

a. Abstract concepts: Do not tag abstract ideas or symbolic references as characters.

Examples: “International cooperation,” “sexual violence,” “<HER>women’s</HER> participation.” In

the latter example, women are the entities that are taking heroic action, while their participation refers to a concept.

b. Entities hoping for/welcoming/thanking/commending something good. These are passive actions.

c. Laws or Treaties.

d. Speakers at the beginning of the input, separated from the content of the speech with a colon.

5. Annotation of character role labels: Only annotate generic role terms (like “victim”) when no other specification is included.

Example: In “Victims of these atrocious crimes have been waiting for justice”, tag “victims of these atrocious crimes” as Victim. However, if specific entities are mentioned, like in “Victims of these atrocious crimes, namely women, have been waiting for justice” only annotate “women” as Victim.

Output Format

Rewrite the entire speech with in-line tags marking the identified characters. Use these tags: Hero: <HER>text</HER> Villain: <VIL>text</VIL> Victim: <VIC>text</VIC>

Combined roles:

Hero and Villain: <HER_VIL>text</HER_VIL>

Hero and Victim: <HER_VIC>text</HER_VIC>

Victim and Villain: <VIC_VIL>text</VIC_VIL>

Hero and Villain and Victim: <HER_VIL_VIC>text</HER_VIL_VIC>

Overlapping annotations should nest tags.

Example: “Those who commit crimes against women, including the peacekeeping personnel, should be brought to book.”

In this case, “women” is a victim within the broader entity “those who commit crimes against women, including the peacekeeping personnel”, which would be entirely annotated as villain.

The output text should be:

“<VIL>Those who commit crimes against <VIC>women</VIC>, including the peacekeeping personnel</VIL>, should be brought to book.”

In certain cases, characters may be portrayed with multiple roles simultaneously. When this happens, annotate the entity with the combined role. Furthermore, the same entity can take on different roles throughout the speech. For instance, while “women” might be classified as <VIC>women</VIC> in one sentence, the same entity can also be classified as <HER>women</HER> in a different part of the speech.

Example:

If the input text is: “The equal right to decision-making and participation, along with women’s empowerment, is crucial to ensure a functioning society and peace and justice in the aftermath of conflicts.”

The output text should be:

“The equal right to decision-making and participation, along with <HER_VIC>women’s</HER_VIC> empowerment, is crucial to ensure a functioning society and peace and justice in the aftermath of conflicts.”

Explanation: In this case, “women” are portrayed both as victims (since they need external help to be empowered) as well as heroes (since they contribute to a peaceful society). When an entity fits into multiple roles based on the context, use the combined tags.

INSTRUCTION: Rewrite the entire speech with the in-line tags included as specified. When annotating, do not rely on any external world knowledge about the entities described in the text. Focus on the way that the speaker portrays the entities in the text. In your output, do not provide any additional explanation or formatting. Make sure to include all spaces from the original text. Multiple spaces are not equivalent to line breaks. Keep the exact same number of spaces between words. Do not replace spaces with new lines. Do not change capitalization, punctuation, or spelling even if you find mistakes. Only include the annotated text—nothing else. The output should strictly be the speech with the appropriate tags for the identified characters.

Annotate this speech:

{ }

Appendix F Few-shot prompt

Your goal: Rewrite the input speech with only in-line character role annotations, using precise tags, leaving spacing and spelling as is, and following all structural and role-based rules.

Task overview You are given a speech from the United Nations Security Council. Your task is to identify and label characters within the speech. Label each identified character as either Hero, Villain, or Victim. These labels can also be combined to mark that an identified character is portrayed as fulfilling several roles, for example, as Hero and Victim. Some character spans may contain other character spans, for example, when a victim entity is mentioned within a villain entity.

For the annotations, follow these annotation guidelines:

Annotation guidelines

1. Character Level Rules:

a. Include entire noun phrases (NP): This includes any numerical or descriptive modifiers. Example: “<HER>60 million Africans</HER> ”

b. Annotate restrictive relative clauses fully. Example: “<HER>parties that perpetuate acts of violence against women and children</HER> ”

c. Include possessive modifiers within the NP:

Example: “<HER>my delegation</HER> ”

d. When multiple characters are listed together as part of the same NP and share a common role or action, tag the entire sequence as one entity. Do not split or annotate each item separately. Example: “Allow me to begin by thanking <HER>the Secretary-General, Mr. Kevin Hyland, Mr. Yury Fedotov and Ms. Ilwad Elman</HER> for their briefings.”

e. Do not include punctuation at the end of annotations. Example: “<VIC>men and boys</VIC>.” f. Do not annotate predicate nominatives or any other descriptions of characters.

Example: “<VIC>Women</VIC> are not only victims.” Here, victims should not be annotated.

Example: “<HER>women</HER> as agents of peace”

g. Do not annotate positions or groups when they are only mentioned in the abstract, without references to the achievements of the particular entity.

Example: “the establishment of the post of the Special Representative of the Secretary General on sexual violence in situations of armed conflict”

2. Role-Definitions:

a. Hero: An entity is labeled as a Hero if it: helps victims and/or defeats villains increases agreement within groups or commitment to a cause recognizes injustice and actively fights to resolve it has the potential to save others, given their abilities, knowledge, positionality/perspective, but who is not able to do so because of structural discrimination makes a call to action calls out or recognizes unjust treatment/violations

b. Victim: An entity is labeled as a Victim if it: is weak, good, or innocent and in need of protection suffers from injustice (e. g., sexual violence, displacement, physical harm) is excluded from decision-making or denied recognition/power does not have the same rights as others

c. Villain: An entity is labeled as a Villain if it: causes anxiety and fear causes people to lose their daily routines causes harm or disrupts peace acts in a way that prevents equal rights and justice for victims has a negative moral reputation

3. Entities to tag: You should only tag:

a. People

b. Organisations

c. Countries

d. Groups

e. UN Resolutions. Only annotate United Nations Resolutions if they are explicitly personified.

Example: In “Resolution 1325 calls for action,” the resolution should be tagged as a hero. In “working towards the implementation of resolution 1325,” the resolution should not be marked.

4. Entities NOT to tag:

a. Abstract concepts: Do not tag abstract ideas or symbolic references as characters.

Examples: “International cooperation,” “sexual violence,” “<HER>women’s</HER> participation.” In

the latter example, women are the entities that are taking heroic action, while their participation refers to a concept.

b. Entities hoping for/welcoming/thanking/commending something good. These are passive actions.

c. Laws or Treaties.

d. Speakers at the beginning of the input, separated from the content of the speech with a colon.

5. Annotation of character role labels: Only annotate generic role terms (like “victim”) when no other specification is included.

Example: In “Victims of these atrocious crimes have been waiting for justice”, tag “victims of these atrocious crimes” as Victim. However, if specific entities are mentioned, like in “Victims of these atrocious crimes, namely women, have been waiting for justice” only annotate “women” as Victim.

Output Format

Rewrite the entire speech with in-line tags marking the identified characters. Use these tags: Hero: <HER>text</HER> Villain: <VIL>text</VIL> Victim: <VIC>text</VIC>

Combined roles:

Hero and Villain: <HER_VIL>text</HER_VIL>

Hero and Victim: <HER_VIC>text</HER_VIC>

Victim and Villain: <VIC_VIL>text</VIC_VIL>

Hero and Villain and Victim: <HER_VIL_VIC>text</HER_VIL_VIC>

Overlapping annotations should nest tags.

Example: “Those who commit crimes against women, including the peacekeeping personnel, should be brought to book.”

In this case, “women” is a victim within the broader entity “those who commit crimes against women, including the peacekeeping personnel”, which would be entirely annotated as villain.

The output text should be:

“<VIL>Those who commit crimes against <VIC>women</VIC>, including the peacekeeping personnel</VIL>, should be brought to book.”

In certain cases, characters may be portrayed with multiple roles simultaneously. When this happens, annotate the entity with the combined role. Furthermore, the same entity can take on different roles throughout the speech. For instance, while “women” might be classified as <VIC>women</VIC> in one sentence, the same entity can also be classified as <HER>women</HER> in a different part of the speech.

Example:

If the input text is: “The equal right to decision-making and participation, along with women’s empowerment, is crucial to ensure a functioning society and peace and justice in the aftermath of conflicts.”

The output text should be:

“The equal right to decision-making and participation, along with <HER_VIC>women’s</HER_VIC> empowerment, is crucial to ensure a functioning society and peace and justice in the aftermath of conflicts.”

Explanation: In this case, “women” are portrayed both as victims (since they need external help to be empowered) as well as heroes (since they contribute to a peaceful society). When an entity fits into multiple roles based on the context, use the combined tags.

Below is an annotated example of how to identify and label characters within a speech from the United Nations Security Council. Use this format to annotate the speech provided.

Annotated Example 1:

Ms. Rossignol (France) (spoke in French): “Allow me to begin by thanking <HER>the Secretary-General, Mr. Kevin Hyland, Mr. Yury Fedotov and Ms. Ilwad Elman</HER> for their briefings. France associates itself with the statement to be delivered by the observer of the European Union. The actions committed by <VIL>Da’esh</VIL> in the Middle East and by <VIL>Boko Haram</VIL> in Africa are a dramatic illustration of the links that exist today between threats to international peace and security and human trafficking. At the global level, trafficking in <VIC>persons</VIC> and trafficking in drugs and counterfeit currency are among the most profitable. Their annual profits are estimated at \$32 billion. It is one of the most extensive forms of trafficking. Sexual exploitation, forced labour, servitude, kidnapping for forced prostitution, rape - the list of atrocities committed in armed conflict is, unfortunately, long. Deriving profits from <VIC>human beings</VIC> and considering <VIC>them</VIC> as merchandise,

<VIL>the traffickers, as well as consumers and users and the clients of sexual exploitation</VIL>, clearly and brutally violate their <VIL>victims'</VIL> human rights and further stoke the causes of conflict. <HER>The international community</HER> has invested heavily in this problem since the beginning of this century, but further efforts are needed in order to address the scourge of human trafficking. I therefore welcome the initiative of <HER>the United Kingdom</HER> during its presidency of the Security Council and thank <HER>that country</HER> for giving us this opportunity to have an exchange on this very important subject. For <HER>France</HER>, the issue of trafficking in <VIC>human beings</VIC> and slavery, especially that of <VIC>women and children</VIC>, is of major importance. The statistics are, unfortunately, well known, but we must constantly point them out: 80 per cent of the victims of trafficking are <VIC>women and children</VIC>. The challenges are also well known: the identification of <VIC>victims</VIC> is still in its infancy, and organized mechanisms for fighting this scourge vary greatly between countries. Despite progress since the entry into force of the Palermo Protocol to Prevent, Suppress and Punish Trafficking in Persons, Especially Women and Children, still too few prosecutions have been initiated in cases involving the crime of human trafficking. <VIC>The victims</VIC> themselves do not always assert their rights and very often are insufficiently protected. In that context, international cooperation must be stepped up so as to increase the geographic coverage of the legislation providing effective protection against <VIL>networks</VIL> and to improve international cooperation aimed at dismantling <VIL>those networks</VIL>. Prevention, protection and the fight against impunity are the three priorities of French diplomacy in the fight against trafficking in <VIC>human beings</VIC>. Since human trafficking is now an integral part of the strategy of <VIL>certain terrorist groups</VIL> and it fuels transnational organized crime, the Security Council has a special responsibility in combating this scourge. The adoption of resolution 2331 (2016), last December, at the initiative of <HER>Spain</HER>, was a major step forward towards better addressing the link between trafficking in <VIC>human beings</VIC>, sexual violence and terrorism. France very much looks forward to the report to be prepared by the Secretary-General by the end of the year. <HER>We</HER> have in place a robust international legal framework and appropriate tools, in particular the United Nations Convention against Transnational Organized Crime and its Protocols and <HER>the United Nations Office on Drugs and Crime</HER>, which is doing sterling work in this field. On International Women's Rights Day, <HER>the President of the French Republic</HER> also announced that <HER>France</HER> would propose an additional protocol to the Convention for the Elimination of All Forms of Discrimination against Women. That protocol would address violence against <VIC>women</VIC> in order to complement the existing international framework. But we must ensure that the obligations arising from that legal framework are effectively implemented. Our words must now be translated into action. Rest assured, Mr. President, that <HER>France</HER> will continue to play its full part in those efforts."

Annotated Example 2: Mr. Juwayeyi (Malawi): <VIC>Those of us who are not in the Security Council</VIC> do not get an opportunity to congratulate a delegation for assuming the presidency of the Council, so it gives me particular pleasure this morning to congratulate <HER>you, Mr. President</HER>, and I am most grateful that during your presidency <HER>you</HER> have taken the initiative to hold this open session on women and peace and security. <HER>My Government</HER> attaches great importance to the protection and security of <VIC>women and girls</VIC>, both in situations of armed conflict and in peace. Wars and armed conflict bring untold suffering and misery to <VIC>communities and nations</VIC>, for they entail devastating and horrific levels of violence and brutality, employing any possible means. Today's wars and conflicts make little distinction between <VIC>soldiers and civilians</VIC> and between <VIC>adults and children</VIC>. Currently, most of the wars and conflicts take place in developing countries, where most of the population lives in rural areas. Often, these conflicts are within countries, rather than across borders. <VIC>Women and children</VIC> constitute a disproportionate number of the affected populations and, therefore, suffer the brunt of violence and brutality. Armed conflict affects <VIC>women and girls</VIC> differently from <VIC>men and boys</VIC>. During armed conflict, not only are <VIC>women and girls</VIC> killed, maimed, abducted, separated from their loved ones, subjected to starvation, malnutrition and forced displacement, but <VIC>they</VIC> are also continually threatened with rape, domestic violence, sexual exploitation and slavery, trafficking, sexual humiliation and mutilation. Rape and sexual violence perpetrated by <VIL>the armed forces,

whether governmental or other actors, including in some instances peacekeeping personnel</VIL>, increases the potential for spreading HIV/AIDS and other sexually transmitted diseases. No wonder most of the HIV/AIDS victims in the developing countries are <VIC>women and girls</VIC>. HIV/AIDS leaves <VIC>millions of children</VIC> orphaned and, in most cases, the responsibility to care for <VIC>them</VIC> rests largely on the shoulders of <HER_VIC>older people</HER_VIC>. All of these harmful and widespread threats to <VIC>women and girls</VIC> have long-term consequences for durable peace, security and development. The sad thing is that in most instances <VIC>the women</VIC> do not know why the wars and armed conflicts erupt, owing to the fact that <VIC>they</VIC> are either under-represented or not represented at all at the decision-making levels. My Government applauds and thanks <HER>the Secretary-General, the United Nations bodies and agencies, non-governmental organizations, international agencies and donor countries</HER> for the efforts they have made to protect and secure peace and security for <VIC>women and girls</VIC>. Various international legal instruments, particularly the Convention on the Elimination of All Forms of Discrimination against Women, the Convention on the Rights of the Child and the corresponding Optional Protocols, have been adopted by <HER>the General Assembly</HER>. <HER>International Criminal Tribunals</HER> have been established for the former Yugoslavia and for Rwanda, and these have made great strides to help end impunity for crimes against <VIC>women and girls</VIC>. Forms of sexual violence are now included as a war crime in the Rome Statute of the International Criminal Court. Beyond its emergency relief responsibilities, <HER>the United Nations</HER> has increasingly become involved in efforts aimed at peacekeeping, peace-making and peace-building. It is gratifying to note that <HER_VIL>the Security Council</HER_VIL>, even though it has taken five decades to do so, has now recognized the importance of <HER_VIC>women's</HER_VIC> role and of their increased participation in the prevention and resolution of conflicts and in peace-building. However, there is still a lot more that needs to be done. Appropriate solutions cannot be achieved if <VIC>women</VIC> are left out of the decision-making machinery. <VIL>You</VIL> are aware, <VIL>Mr. President</VIL>, that <VIC>women</VIC> continue to be under-represented in all peacekeeping, peace-making and peace-building efforts, including in the Department of Peacekeeping Operations in the United Nations. This should not be allowed to continue. In the long term however, the only way to truly ensure the protection and security of <VIC>women and girls</VIC> is to prevent wars and armed conflicts from taking place. Major root causes of most of the recent wars and armed conflicts have included poverty and lack of respect for human rights. These ills must be addressed first. My delegation recalls that, at the end of the Millennium Summit, world leaders pledged to eradicate poverty and make the right to development a reality for everyone. This means promoting equality between men and women in decision-making. This further means the involvement and full participation of <HER_VIC>women</HER_VIC> in all issues, including peacekeeping, peacemaking and peace-building, as well as at the negotiating table, from the grassroots level to the decision-making levels. <HER>My delegation</HER> requests the Secretary-General and the Security Council to urge Member States to ensure that training in human rights and peacekeeping, peacemaking and peace-building includes everyone - civilians, soldiers, the police, civil society, the women themselves and peacekeeping personnel. <VIL>Those who commit crimes against <VIC>women</VIC>, including the peacekeeping personnel</VIL>, should be brought to book. Let <HER>us</HER> heed the <VIC>women's</VIC> cry for an equal opportunity to voice their ideas in official peace negotiations. And let <HER>us</HER> act now.

INSTRUCTION: Rewrite the entire speech with the in-line tags included as specified. When annotating, do not rely on any external world knowledge about the entities described in the text. Focus on the way that the speaker portrays the entities in the text. In your output, do not provide any additional explanation or formatting. Make sure to include all spaces from the original text. Multiple spaces are not equivalent to line breaks. Keep the exact same number of spaces between words. Do not replace spaces with new lines. Do not change capitalization, punctuation, or spelling even if you find mistakes. Only include the annotated text—nothing else. The output should strictly be the speech with the appropriate tags for the identified characters.

Annotate this speech:

{ }